

```
In [9]: import pandas as pd  
import numpy as np  
import seaborn as sns
```

```
In [10]: df = pd.read_csv('./data/DATA/fake_reg.csv')
```

```
In [11]: df.head()
```

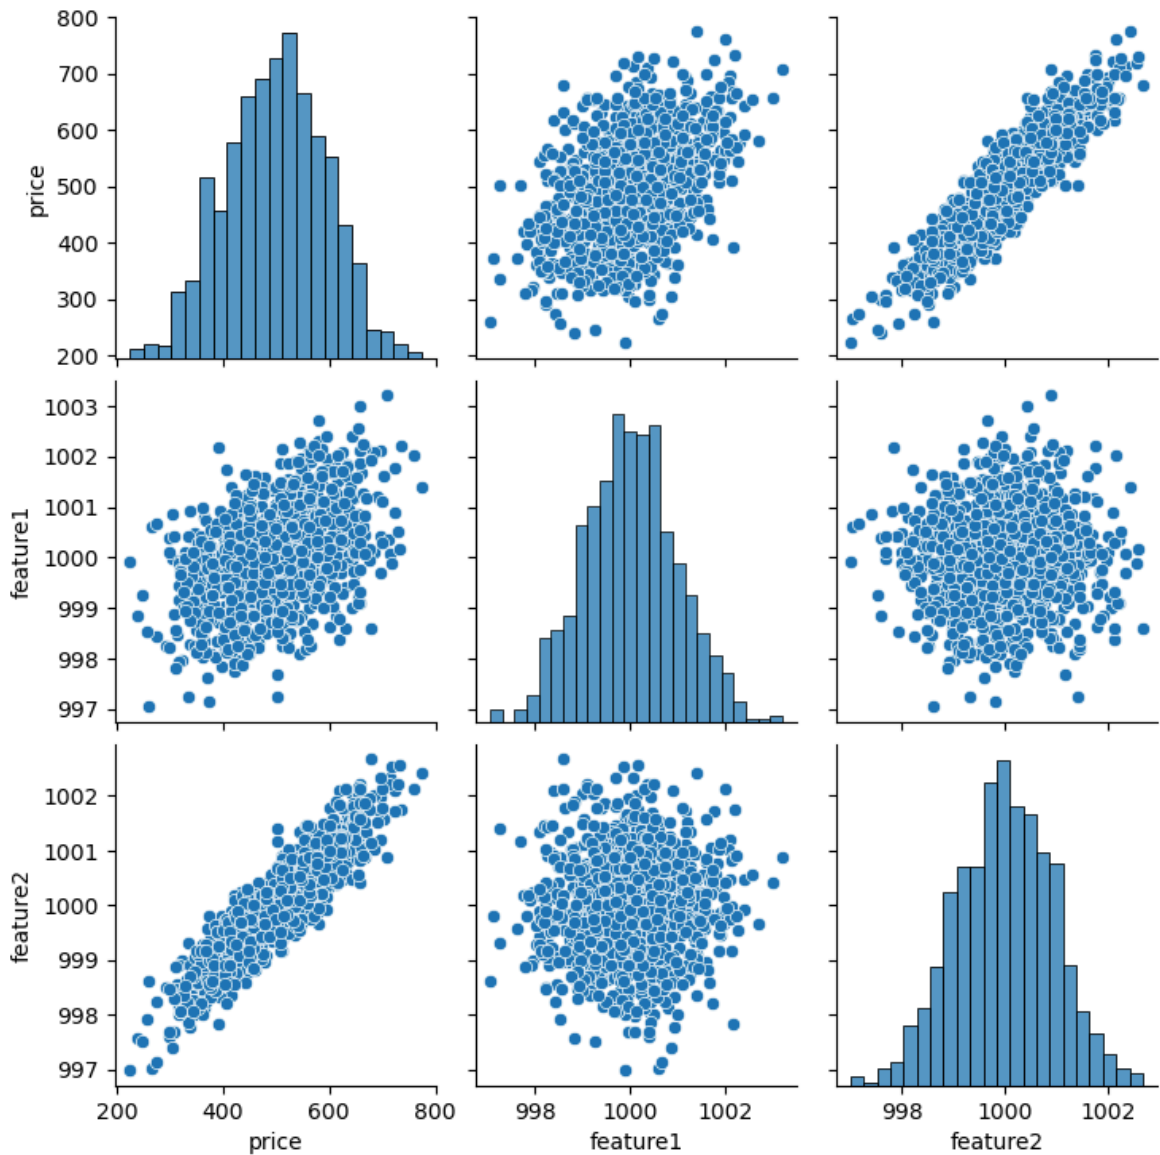
```
Out[11]:
```

	price	feature1	feature2
0	461.527929	999.787558	999.766096
1	548.130011	998.861615	1001.042403
2	410.297162	1000.070267	998.844015
3	540.382220	999.952251	1000.440940
4	546.024553	1000.446011	1000.338531

```
In [12]: sns.pairplot(df)
```

```
C:\learnings\envs\deeplearning\lib\site-packages\seaborn\axisgrid.py:123: UserWarning: The figure layout has changed to tight  
self._figure.tight_layout(*args, **kwargs)
```

```
Out[12]: <seaborn.axisgrid.PairGrid at 0x27b29bacfa0>
```



```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: X = df[['feature1', 'feature2']].values
```

```
In [16]: y = df['price'].values
```

```
In [18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
```

```
In [19]: X_train.shape
```

```
Out[19]: (700, 2)
```

```
In [20]: X_test.shape
```

```
Out[20]: (300, 2)
```

```
In [21]: from sklearn.preprocessing import MinMaxScaler
```

```
In [25]: #help(MinMaxScaler)
```

```
In [22]: scaler = MinMaxScaler()
```

```
In [23]: scaler.fit(X_train)
```

```
Out[23]: ▼ MinMaxScaler  
MinMaxScaler()
```

```
In [24]: X_train = scaler.transform(X_train)
```

```
In [25]: X_test = scaler.transform(X_test)
```

```
In [26]: X_train.max()
```

```
Out[26]: 1.0
```

```
In [28]: from tensorflow.keras.models import Sequential
```

```
In [29]: from tensorflow.keras.layers import Dense
```

```
In [30]: help(Sequential)
```

Help on class Sequential in module keras.engine.sequential:

```
class Sequential(keras.engine.functional.Functional)
| Sequential(*args, **kwargs)
|
| `Sequential` groups a linear stack of layers into a `tf.keras.Model`.
|
| `Sequential` provides training and inference features on this model.
|
| Examples:
|
| ```python
| # Optionally, the first layer can receive an `input_shape` argument:
| model = tf.keras.Sequential()
| model.add(tf.keras.layers.Dense(8, input_shape=(16,)))
| # Afterwards, we do automatic shape inference:
| model.add(tf.keras.layers.Dense(4))
|
| # This is identical to the following:
| model = tf.keras.Sequential()
| model.add(tf.keras.Input(shape=(16,)))
| model.add(tf.keras.layers.Dense(8))
|
| # Note that you can also omit the `input_shape` argument.
| # In that case the model doesn't have any weights until the first call
| # to a training/evaluation method (since it isn't yet built):
| model = tf.keras.Sequential()
| model.add(tf.keras.layers.Dense(8))
| model.add(tf.keras.layers.Dense(4))
| # model.weights not created yet
|
| # Whereas if you specify the input shape, the model gets built
| # continuously as you are adding layers:
| model = tf.keras.Sequential()
| model.add(tf.keras.layers.Dense(8, input_shape=(16,)))
| model.add(tf.keras.layers.Dense(4))
| len(model.weights)
| # Returns "4"
|
| # When using the delayed-build pattern (no input shape specified), you can
| # choose to manually build your model by calling
| # `build(batch_input_shape)`:
| model = tf.keras.Sequential()
| model.add(tf.keras.layers.Dense(8))
| model.add(tf.keras.layers.Dense(4))
| model.build((None, 16))
| len(model.weights)
| # Returns "4"
|
| # Note that when using the delayed-build pattern (no input shape specified),
| # the model gets built the first time you call `fit`, `eval`, or `predict`,
| # or the first time you call the model on some input data.
| model = tf.keras.Sequential()
| model.add(tf.keras.layers.Dense(8))
| model.add(tf.keras.layers.Dense(1))
| model.compile(optimizer='sgd', loss='mse')
| # This builds the model for the first time:
| model.fit(x, y, batch_size=32, epochs=10)
| ```
```

Method resolution order:

```
Sequential
keras.engine.functional.Functional
keras.engine.training.Model
keras.engine.base_layer.Layer
tensorflow.python.module.module.Module
tensorflow.python.trackable.autotrackable.AutoTrackable
tensorflow.python.trackable.base.Trackable
keras.utils.version_utils.LayerVersionSelector
keras.utils.version_utils.ModelVersionSelector
builtins.object
```

Methods defined here:

`__init__(self, layers=None, name=None)`
Creates a `Sequential` model instance.

Args:

layers: Optional list of layers to add to the model.
name: Optional name for the model.

`add(self, layer)`
Adds a layer instance on top of the layer stack.

Args:

layer: layer instance.

Raises:

TypeError: If `layer` is not a layer instance.
ValueError: In case the `layer` argument does not know its input shape.
ValueError: In case the `layer` argument has multiple output tensors, or is already connected somewhere else (forbidden in `Sequential` models).

`build(self, input_shape=None)`
Builds the model based on input shapes received.

This is to be used for subclassed models, which do not know at instantiation time what their inputs look like.

This method only exists for users who want to call `model.build()` in a standalone way (as a substitute for calling the model on real data to build it). It will never be called by the framework (and thus it will never throw unexpected errors in an unrelated workflow).

Args:

input_shape: Single tuple, `TensorShape` instance, or list/dict of shapes, where shapes are tuples, integers, or `TensorShape` instances.

Raises:

ValueError:

1. In case of invalid user-provided data (not of type tuple, list, `TensorShape`, or dict).
2. If the model requires call arguments that are agnostic to the input shapes (positional or keyword arg in call signature).
3. If not all layers were properly built.
4. If float type inputs are not supported within the layers.

In each of these cases, the user should build their model by calling it on real tensor data.

```
call(self, inputs, training=None, mask=None)
```

Calls the model on new inputs.

In this case `call` just reapplies all ops in the graph to the new inputs (e.g. build a new computational graph from the provided inputs).

Args:

`inputs`: A tensor or list of tensors.
`training`: Boolean or boolean scalar tensor, indicating whether to run the `Network` in training mode or inference mode.
`mask`: A mask or list of masks. A mask can be either a tensor or None (no mask).

Returns:

A tensor if there is a single output, or a list of tensors if there are more than one outputs.

```
compute_mask(self, inputs, mask)
```

Computes an output mask tensor.

Args:

`inputs`: Tensor or list of tensors.
`mask`: Tensor or list of tensors.

Returns:

None or a tensor (or list of tensors, one per output tensor of the layer).

```
compute_output_shape(self, input_shape)
```

Computes the output shape of the layer.

This method will cause the layer's state to be built, if that has not happened before. This requires that the layer will later be used with inputs that match the input shape provided here.

Args:

`input_shape`: Shape tuple (tuple of integers) or list of shape tuples (one per output tensor of the layer). Shape tuples can include None for free dimensions, instead of an integer.

Returns:

An input shape tuple.

```
get_config(self)
```

Returns the config of the `Model`.

Config is a Python dictionary (serializable) containing the configuration of an object, which in this case is a `Model`. This allows the `Model` to be reinstantiated later (without its trained weights) from this configuration.

Note that `get_config()` does not guarantee to return a fresh copy of dict every time it is called. The callers should make a copy of the returned dict if they want to modify it.

Developers of subclassed `Model` are advised to override this method, and continue to update the dict from `super(MyModel, self).get_config()` to provide the proper configuration of this `Model`. The default config is an empty dict. Optionally, raise `NotImplementedError` to allow Keras to attempt a default serialization.

Returns:

Python dictionary containing the configuration of this `Model`.

`pop(self)`

Removes the last layer in the model.

Raises:

`TypeError`: if there are no layers in the model.

Class methods defined here:

`from_config(config, custom_objects=None)` from `builtins.type`
Creates a layer from its config.

This method is the reverse of `get_config`, capable of instantiating the same layer from the config dictionary. It does not handle layer connectivity (handled by `Network`), nor weights (handled by `set_weights`).

Args:

`config`: A Python dictionary, typically the output of `get_config`.

Returns:

A layer instance.

Readonly properties defined here:

`layers`

Data descriptors defined here:

`input_spec`

`InputSpec` instance(s) describing the input format for this layer.

When you create a layer subclass, you can set `self.input_spec` to enable the layer to run input compatibility checks when it is called. Consider a `Conv2D` layer: it can only be called on a single input tensor of rank 4. As such, you can set, in `__init__`:

```
```python
self.input_spec = tf.keras.layers.InputSpec(ndim=4)
```
```

Now, if you try to call the layer on an input that isn't rank 4 (for instance, an input of shape `(2,)`, it will raise a nicely-formatted error:

```
```
```

`ValueError: Input 0 of layer conv2d is incompatible with the layer:`

```
expected ndim=4, found ndim=1. Full shape received: [2]
```

```

Input checks that can be specified via ``input_spec`` include:

- Structure (e.g. a single input, a list of 2 inputs, etc)
- Shape
- Rank (ndim)
- Dtype

For more information, see ``tf.keras.layers.InputSpec``.

Returns:

A ``tf.keras.layers.InputSpec`` instance, or nested structure thereof.

Methods inherited from `keras.engine.functional.Functional``:

`get_weight_paths(self)`

Retrieve all the variables and their paths for the model.

The variable path (string) is a stable key to indentify a ``tf.Variable`` instance owned by the model. It can be used to specify variable-specific configurations (e.g. `DTensor`, quantization) from a global view.

This method returns a dict with weight object paths as keys and the corresponding ``tf.Variable`` instances as values.

Note that if the model is a subclassed model and the weights haven't been initialized, an empty dict will be returned.

Returns:

A dict where keys are variable paths and values are ``tf.Variable`` instances.

Example:

```
```python
```

```
class SubclassModel(tf.keras.Model):
```

```
 def __init__(self, name=None):
 super().__init__(name=name)
 self.d1 = tf.keras.layers.Dense(10)
 self.d2 = tf.keras.layers.Dense(20)
```

```
 def call(self, inputs):
 x = self.d1(inputs)
 return self.d2(x)
```

```
model = SubclassModel()
model(tf.zeros((10, 10)))
weight_paths = model.get_weight_paths()
weight_paths:
{
'd1.kernel': model.d1.kernel,
'd1.bias': model.d1.bias,
'd2.kernel': model.d2.kernel,
'd2.bias': model.d2.bias,
}
```

```
Functional model
```



```

inputs = tf.keras.Input((10,), batch_size=10)
x = tf.keras.layers.Dense(20, name='d1')(inputs)
output = tf.keras.layers.Dense(30, name='d2')(x)
model = tf.keras.Model(inputs, output)
d1 = model.layers[1]
d2 = model.layers[2]
weight_paths = model.get_weight_paths()
weight_paths:
{
'd1.kernel': d1.kernel,
'd1.bias': d1.bias,
'd2.kernel': d2.kernel,
'd2.bias': d2.bias,
}
```

```

Readonly properties inherited from keras.engine.functional.Functional:

input

Retrieves the input tensor(s) of a layer.

Only applicable if the layer has exactly one input,
i.e. if it is connected to one incoming layer.

Returns:

Input tensor or list of input tensors.

Raises:

RuntimeError: If called in Eager mode.

AttributeError: If no inbound nodes are found.

input_shape

Retrieves the input shape(s) of a layer.

Only applicable if the layer has exactly one input,
i.e. if it is connected to one incoming layer, or if all inputs
have the same shape.

Returns:

Input shape, as an integer shape tuple
(or list of shape tuples, one tuple per input tensor).

Raises:

AttributeError: if the layer has no defined input_shape.

RuntimeError: if called in Eager mode.

output

Retrieves the output tensor(s) of a layer.

Only applicable if the layer has exactly one output,
i.e. if it is connected to one incoming layer.

Returns:

Output tensor or list of output tensors.

Raises:

AttributeError: if the layer is connected to more than one incoming
layers.

RuntimeError: if called in Eager mode.

output_shape

Retrieves the output shape(s) of a layer.

Only applicable if the layer has one output,
or if all outputs have the same shape.

Returns:

Output shape, as an integer shape tuple
(or list of shape tuples, one tuple per output tensor).

Raises:

`AttributeError`: if the layer has no defined output shape.
`RuntimeError`: if called in Eager mode.

Methods inherited from `keras.engine.training.Model`:

`__call__(self, *args, **kwargs)`

`__copy__(self)`

`__deepcopy__(self, memo)`

`__reduce__(self)`
Helper for pickle.

`__setattr__(self, name, value)`
Support `self.foo = trackable syntax`.

`compile(self, optimizer='rmsprop', loss=None, metrics=None, loss_weights=None, weighted_metrics=None, run_eagerly=None, steps_per_execution=None, jit_compile=None, **kwargs)`

Configures the model for training.

Example:

```
```python
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
 loss=tf.keras.losses.BinaryCrossentropy(),
 metrics=[tf.keras.metrics.BinaryAccuracy(),
 tf.keras.metrics.FalseNegatives()])
```
```

Args:

`optimizer`: String (name of optimizer) or optimizer instance. See ``tf.keras.optimizers``.

`loss`: Loss function. May be a string (name of loss function), or a ``tf.keras.losses.Loss`` instance. See ``tf.keras.losses``. A loss function is any callable with the signature ``loss = fn(y_true, y_pred)``, where ``y_true`` are the ground truth values, and ``y_pred`` are the model's predictions.

``y_true`` should have shape

``(batch_size, d0, .. dN)`` (except in the case of sparse loss functions such as

sparse categorical crossentropy which expects integer arrays of shape ``(batch_size, d0, .. dN-1)``).

``y_pred`` should have shape ``(batch_size, d0, .. dN)``.

The loss function should return a float tensor.

If a custom ``Loss`` instance is

used and reduction is set to `None`, return value has shape `(batch_size, d0, .. dN-1)` i.e. per-sample or per-timestep loss values; otherwise, it is a scalar. If the model has multiple outputs, you can use a different loss on each output by passing a dictionary or a list of losses. The loss value that will be minimized by the model will then be the sum of all individual losses, unless `loss_weights` is specified.

metrics: List of metrics to be evaluated by the model during training and testing. Each of this can be a string (name of a built-in function), function or a `tf.keras.metrics.Metric` instance. See `tf.keras.metrics`. Typically you will use `metrics=['accuracy']`.

A function is any callable with the signature `result = fn(y_true, y_pred)`. To specify different metrics for different outputs of a multi-output model, you could also pass a dictionary, such as `metrics={'output_a': 'accuracy', 'output_b': ['accuracy', 'mse']}`. You can also pass a list to specify a metric or a list of metrics for each output, such as `metrics=[['accuracy'], ['accuracy', 'mse']]` or `metrics=['accuracy', ['accuracy', 'mse']]`. When you pass the strings 'accuracy' or 'acc', we convert this to one of `tf.keras.metrics.BinaryAccuracy`, `tf.keras.metrics.CategoricalAccuracy`, `tf.keras.metrics.SparseCategoricalAccuracy` based on the loss function used and the model output shape. We do a similar conversion for the strings 'crossentropy' and 'ce' as well. The metrics passed here are evaluated without sample weighting; if you would like sample weighting to apply, you can specify your metrics via the `weighted_metrics` argument instead.

loss_weights: Optional list or dictionary specifying scalar coefficients (Python floats) to weight the loss contributions of different model outputs. The loss value that will be minimized by the model will then be the *weighted sum* of all individual losses, weighted by the `loss_weights` coefficients. If a list, it is expected to have a 1:1 mapping to the model's outputs. If a dict, it is expected to map output names (strings) to scalar coefficients.

weighted_metrics: List of metrics to be evaluated and weighted by `sample_weight` or `class_weight` during training and testing.

run_eagerly: Bool. Defaults to `False`. If `True`, this `Model`'s logic will not be wrapped in a `tf.function`. Recommended to leave this as `None` unless your `Model` cannot be run inside a `tf.function`. `run_eagerly=True` is not supported when using `tf.distribute.experimental.ParameterServerStrategy`.

steps_per_execution: Int. Defaults to 1. The number of batches to run during each `tf.function` call. Running multiple batches inside a single `tf.function` call can greatly improve performance on TPUs or small models with a large Python overhead. At most, one full epoch will be run each execution. If a number larger than the size of the epoch is passed, the execution will be truncated to the size of the epoch. Note that if `steps_per_execution` is set to `N`, `Callback.on_batch_begin` and `Callback.on_batch_end` methods will only be called every `N` batches (i.e. before/after each `tf.function` execution).

jit_compile: If `True`, compile the model training step with XLA. [XLA](https://www.tensorflow.org/xla) is an optimizing compiler for machine learning. `jit_compile` is not enabled for by default. This option cannot be enabled with `run_eagerly=True`. Note that `jit_compile=True`

may not necessarily work for all models.
 For more information on supported operations please refer to the
[\[XLA documentation\]\(https://www.tensorflow.org/xla\)](https://www.tensorflow.org/xla).
 Also refer to
[\[known XLA issues\]\(https://www.tensorflow.org/xla/known_issues\)](https://www.tensorflow.org/xla/known_issues)
 for more details.
 **kwargs: Arguments supported for backwards compatibility only.

`compute_loss(self, x=None, y=None, y_pred=None, sample_weight=None)`
 Compute the total loss, validate it, and return it.

Subclasses can optionally override this method to provide custom loss computation logic.

Example:

```
```python
class MyModel(tf.keras.Model):

 def __init__(self, *args, **kwargs):
 super(MyModel, self).__init__(*args, **kwargs)
 self.loss_tracker = tf.keras.metrics.Mean(name='loss')

 def compute_loss(self, x, y, y_pred, sample_weight):
 loss = tf.reduce_mean(tf.math.squared_difference(y_pred, y))
 loss += tf.add_n(self.losses)
 self.loss_tracker.update_state(loss)
 return loss

 def reset_metrics(self):
 self.loss_tracker.reset_states()

 @property
 def metrics(self):
 return [self.loss_tracker]

tensors = tf.random.uniform((10, 10)), tf.random.uniform((10,))
dataset = tf.data.Dataset.from_tensor_slices(tensors).repeat().batch(1)

inputs = tf.keras.layers.Input(shape=(10,), name='my_input')
outputs = tf.keras.layers.Dense(10)(inputs)
model = MyModel(inputs, outputs)
model.add_loss(tf.reduce_sum(outputs))

optimizer = tf.keras.optimizers.SGD()
model.compile(optimizer, loss='mse', steps_per_execution=10)
model.fit(dataset, epochs=2, steps_per_epoch=10)
print('My custom loss: ', model.loss_tracker.result().numpy())
```
```

Args:

x: Input data.
 y: Target data.
 y_pred: Predictions returned by the model (output of `model(x)`)
 sample_weight: Sample weights for weighting the loss function.

Returns:

The total loss as a `tf.Tensor`, or `None` if no loss results (which is the case when called by `Model.test_step`).

`compute_metrics(self, x, y, y_pred, sample_weight)`

Update metric states and collect all metrics to be returned.

Subclasses can optionally override this method to provide custom metric updating and collection logic.

Example:

```
```python
class MyModel(tf.keras.Sequential):

 def compute_metrics(self, x, y, y_pred, sample_weight):

 # This super call updates `self.compiled_metrics` and returns
 # results for all metrics listed in `self.metrics`.
 metric_results = super(MyModel, self).compute_metrics(
 x, y, y_pred, sample_weight)

 # Note that `self.custom_metric` is not listed in `self.metrics`.
 self.custom_metric.update_state(x, y, y_pred, sample_weight)
 metric_results['custom_metric_name'] = self.custom_metric.result()
 return metric_results
```
```

Args:

x: Input data.
y: Target data.
y_pred: Predictions returned by the model (output of `model.call(x)`)
sample_weight: Sample weights for weighting the loss function.

Returns:

A `dict` containing values that will be passed to
`tf.keras.callbacks.CallbackList.on_train_batch_end()`. Typically, the
values of the metrics listed in `self.metrics` are returned. Example:
`{'loss': 0.2, 'accuracy': 0.7}`.

evaluate(self, x=None, y=None, batch_size=None, verbose='auto', sample_weight=None, steps=None, callbacks=None, max_queue_size=10, workers=1, use_multiprocessing=False, return_dict=False, **kwargs)

Returns the loss value & metrics values for the model in test mode.

Computation is done in batches (see the `batch_size` arg.)

Args:

x: Input data. It could be:

- A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
- A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).
- A dict mapping input names to the corresponding array/tensors, if the model has named inputs.
- A `tf.data` dataset. Should return a tuple of either `(inputs, targets)` or `(inputs, targets, sample_weights)`.
- A generator or `keras.utils.Sequence` returning `(inputs, targets)` or `(inputs, targets, sample_weights)`.

A more detailed description of unpacking behavior for iterator types (Dataset, generator, Sequence) is given in the `Unpacking behavior for iterator-like inputs` section of `Model.fit`.

y: Target data. Like the input data `x`, it could be either Numpy array(s) or TensorFlow tensor(s). It should be consistent with `x` (you cannot have Numpy inputs and tensor targets, or inversely).

If `x` is a dataset, generator or `keras.utils.Sequence` instance, `y` should not be specified (since targets will be obtained from the iterator/dataset).

`batch_size`: Integer or `None`. Number of samples per batch of computation. If unspecified, `batch_size` will default to 32. Do not specify the `batch_size` if your data is in the form of a dataset, generators, or `keras.utils.Sequence` instances (since they generate batches).

`verbose`: `"auto"`, 0, 1, or 2. Verbosity mode.
 0 = silent, 1 = progress bar, 2 = single line.
`"auto"` defaults to 1 for most cases, and to 2 when used with `ParameterServerStrategy`. Note that the progress bar is not particularly useful when logged to a file, so `verbose=2` is recommended when not running interactively (e.g. in a production environment).

`sample_weight`: Optional Numpy array of weights for the test samples, used for weighting the loss function. You can either pass a flat (1D) Numpy array with the same length as the input samples (1:1 mapping between weights and samples), or in the case of temporal data, you can pass a 2D array with shape `(samples, sequence_length)`, to apply a different weight to every timestep of every sample. This argument is not supported when `x` is a dataset, instead pass sample weights as the third element of `x`.

`steps`: Integer or `None`. Total number of steps (batches of samples) before declaring the evaluation round finished. Ignored with the default value of `None`. If `x` is a `tf.data` dataset and `steps` is `None`, 'evaluate' will run until the dataset is exhausted. This argument is not supported with array inputs.

`callbacks`: List of `keras.callbacks.Callback` instances. List of callbacks to apply during evaluation. See [\[callbacks\]\(/api_docs/python/tf/keras/callbacks\)](#).

`max_queue_size`: Integer. Used for generator or `keras.utils.Sequence` input only. Maximum size for the generator queue. If unspecified, `max_queue_size` will default to 10.

`workers`: Integer. Used for generator or `keras.utils.Sequence` input only. Maximum number of processes to spin up when using process-based threading. If unspecified, `workers` will default to 1.

`use_multiprocessing`: Boolean. Used for generator or `keras.utils.Sequence` input only. If `True`, use process-based threading. If unspecified, `use_multiprocessing` will default to `False`. Note that because this implementation relies on multiprocessing, you should not pass non-picklable arguments to the generator as they can't be passed easily to children processes.

`return_dict`: If `True`, loss and metric results are returned as a dict, with each key being the name of the metric. If `False`, they are returned as a list.

`**kwargs`: Unused at this time.

See the discussion of 'Unpacking behavior for iterator-like inputs' for `Model.fit`.

Returns:

Scalar test loss (if the model has a single output and no metrics) or list of scalars (if the model has multiple outputs and/or metrics). The attribute `model.metrics_names` will give you the display labels for the scalar outputs.

```

    Raises:
        RuntimeError: If `model.evaluate` is wrapped in a `tf.function`.

    evaluate_generator(self, generator, steps=None, callbacks=None, max_queue_size=10, workers=1, use_multiprocessing=False, verbose=0)
        Evaluates the model on a data generator.

    DEPRECATED:
        `Model.evaluate` now supports generators, so there is no longer any need to use this endpoint.

    fit(self, x=None, y=None, batch_size=None, epochs=1, verbose='auto', callbacks=None, validation_split=0.0, validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None, validation_steps=None, validation_batch_size=None, validation_freq=1, max_queue_size=10, workers=1, use_multiprocessing=False)
        Trains the model for a fixed number of epochs (iterations on a dataset).

    Args:
        x: Input data. It could be:
            - A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
            - A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).
            - A dict mapping input names to the corresponding array/tensors, if the model has named inputs.
            - A `tf.data` dataset. Should return a tuple of either `(inputs, targets)` or `(inputs, targets, sample_weights)`.
            - A generator or `keras.utils.Sequence` returning `(inputs, targets)` or `(inputs, targets, sample_weights)`.
            - A `tf.keras.utils.experimental.DatasetCreator`, which wraps a callable that takes a single argument of type `tf.distribute.InputContext`, and returns a `tf.data.Dataset`. `DatasetCreator` should be used when users prefer to specify the per-replica batching and sharding logic for the `Dataset`. See `tf.keras.utils.experimental.DatasetCreator` doc for more information.
            A more detailed description of unpacking behavior for iterator types (Dataset, generator, Sequence) is given below. If these include `sample_weights` as a third component, note that sample weighting applies to the `weighted_metrics` argument but not the `metrics` argument in `compile()`. If using `tf.distribute.experimental.ParameterServerStrategy`, only `DatasetCreator` type is supported for `x`.
        y: Target data. Like the input data `x`, it could be either Numpy array(s) or TensorFlow tensor(s). It should be consistent with `x` (you cannot have Numpy inputs and tensor targets, or inversely). If `x` is a dataset, generator, or `keras.utils.Sequence` instance, `y` should not be specified (since targets will be obtained from `x`).
        batch_size: Integer or `None`.
            Number of samples per gradient update.
            If unspecified, `batch_size` will default to 32.
            Do not specify the `batch_size` if your data is in the form of datasets, generators, or `keras.utils.Sequence` instances (since they generate batches).
        epochs: Integer. Number of epochs to train the model.
            An epoch is an iteration over the entire `x` and `y` data provided

```

(unless the ``steps_per_epoch`` flag is set to something other than `None`).

Note that in conjunction with ``initial_epoch``, ``epochs`` is to be understood as "final epoch".

The model is not trained for a number of iterations given by ``epochs``, but merely until the epoch of index ``epochs`` is reached.

`verbose`: 'auto', 0, 1, or 2. Verbosity mode.

- 0 = silent, 1 = progress bar, 2 = one line per epoch.
- 'auto' defaults to 1 for most cases, but 2 when used with ``ParameterServerStrategy``. Note that the progress bar is not particularly useful when logged to a file, so `verbose=2` is recommended when not running interactively (eg, in a production environment).

`callbacks`: List of ``keras.callbacks.Callback`` instances.

List of callbacks to apply during training.

See ``tf.keras.callbacks``. Note

- ``tf.keras.callbacks.ProgbarLogger`` and ``tf.keras.callbacks.History`` callbacks are created automatically and need not be passed into ``model.fit``.
- ``tf.keras.callbacks.ProgbarLogger`` is created or not based on ``verbose`` argument to ``model.fit``.

Callbacks with batch-level calls are currently unsupported with ``tf.distribute.experimental.ParameterServerStrategy``, and users are advised to implement epoch-level calls instead with an appropriate ``steps_per_epoch`` value.

`validation_split`: Float between 0 and 1.

Fraction of the training data to be used as validation data.

The model will set apart this fraction of the training data, will not train on it, and will evaluate the loss and any model metrics on this data at the end of each epoch.

The validation data is selected from the last samples in the ``x`` and ``y`` data provided, before shuffling. This argument is not supported when ``x`` is a dataset, generator or ``keras.utils.Sequence`` instance.

If both ``validation_data`` and ``validation_split`` are provided, ``validation_data`` will override ``validation_split``.

``validation_split`` is not yet supported with ``tf.distribute.experimental.ParameterServerStrategy``.

`validation_data`: Data on which to evaluate the loss and any model metrics at the end of each epoch.

The model will not be trained on this data. Thus, note the fact that the validation loss of data provided using ``validation_split`` or ``validation_data`` is not affected by regularization layers like noise and dropout.

``validation_data`` will override ``validation_split``.

``validation_data`` could be:

- A tuple ``(x_val, y_val)`` of Numpy arrays or tensors.
- A tuple ``(x_val, y_val, val_sample_weights)`` of NumPy arrays.
- A ``tf.data.Dataset``.
- A Python generator or ``keras.utils.Sequence`` returning ``(inputs, targets)`` or ``(inputs, targets, sample_weights)``.

``validation_data`` is not yet supported with ``tf.distribute.experimental.ParameterServerStrategy``.

`shuffle`: Boolean (whether to shuffle the training data before each epoch) or str (for 'batch'). This argument is ignored when ``x`` is a generator or an object of `tf.data.Dataset`.

'batch' is a special option for dealing

with the limitations of HDF5 data; it shuffles in batch-sized chunks. Has no effect when `steps_per_epoch` is not `None`.

class_weight: Optional dictionary mapping class indices (integers) to a weight (float) value, used for weighting the loss function (during training only). This can be useful to tell the model to "pay more attention" to samples from an under-represented class.

sample_weight: Optional Numpy array of weights for the training samples, used for weighting the loss function (during training only). You can either pass a flat (1D) Numpy array with the same length as the input samples (1:1 mapping between weights and samples), or in the case of temporal data, you can pass a 2D array with shape `(samples, sequence_length)`, to apply a different weight to every timestep of every sample. This argument is not supported when `x` is a dataset, generator, or `keras.utils.Sequence` instance, instead provide the `sample_weights` as the third element of `x`. Note that sample weighting does not apply to metrics specified via the `metrics` argument in `compile()`. To apply sample weighting to your metrics, you can specify them via the `weighted_metrics` in `compile()` instead.

initial_epoch: Integer. Epoch at which to start training (useful for resuming a previous training run).

steps_per_epoch: Integer or `None`. Total number of steps (batches of samples) before declaring one epoch finished and starting the next epoch. When training with input tensors such as TensorFlow data tensors, the default `None` is equal to the number of samples in your dataset divided by the batch size, or 1 if that cannot be determined. If `x` is a `tf.data` dataset, and `'steps_per_epoch'` is `None`, the epoch will run until the input dataset is exhausted. When passing an infinitely repeating dataset, you must specify the `steps_per_epoch` argument. If `'steps_per_epoch=-1'` the training will run indefinitely with an infinitely repeating dataset. This argument is not supported with array inputs. When using `tf.distribute.experimental.ParameterServerStrategy`: * `'steps_per_epoch=None'` is not supported.

validation_steps: Only relevant if `'validation_data'` is provided and is a `tf.data` dataset. Total number of steps (batches of samples) to draw before stopping when performing validation at the end of every epoch. If `'validation_steps'` is `None`, validation will run until the `'validation_data'` dataset is exhausted. In the case of an infinitely repeated dataset, it will run into an infinite loop. If `'validation_steps'` is specified and only part of the dataset will be consumed, the evaluation will start from the beginning of the dataset at each epoch. This ensures that the same validation samples are used every time.

validation_batch_size: Integer or `None`. Number of samples per validation batch. If unspecified, will default to `'batch_size'`. Do not specify the `'validation_batch_size'` if your data is in the form of datasets, generators, or `keras.utils.Sequence` instances (since they generate batches).

`validation_freq`: Only relevant if validation data is provided. Integer or ``collections.abc.Container`` instance (e.g. list, tuple, etc.). If an integer, specifies how many training epochs to run before a new validation run is performed, e.g. ``validation_freq=2`` runs validation every 2 epochs. If a Container, specifies the epochs on which to run validation, e.g. ``validation_freq=[1, 2, 10]`` runs validation at the end of the 1st, 2nd, and 10th epochs.

`max_queue_size`: Integer. Used for generator or ``keras.utils.Sequence`` input only. Maximum size for the generator queue. If unspecified, ``max_queue_size`` will default to 10.

`workers`: Integer. Used for generator or ``keras.utils.Sequence`` input only. Maximum number of processes to spin up when using process-based threading. If unspecified, ``workers`` will default to 1.

`use_multiprocessing`: Boolean. Used for generator or ``keras.utils.Sequence`` input only. If ``True``, use process-based threading. If unspecified, ``use_multiprocessing`` will default to ``False``. Note that because this implementation relies on multiprocessing, you should not pass non-picklable arguments to the generator as they can't be passed easily to children processes.

Unpacking behavior for iterator-like inputs:

A common pattern is to pass a `tf.data.Dataset`, generator, or `tf.keras.utils.Sequence` to the ``x`` argument of `fit`, which will in fact yield not only features (`x`) but optionally targets (`y`) and sample weights. Keras requires that the output of such iterator-like be unambiguous. The iterator should return a tuple of length 1, 2, or 3, where the optional second and third elements will be used for `y` and `sample_weight` respectively. Any other type provided will be wrapped in a length one tuple, effectively treating everything as ``x``. When yielding dicts, they should still adhere to the top-level tuple structure.

e.g. ``({`x0`: x0, `x1`: x1}, y)``. Keras will not attempt to separate features, targets, and weights from the keys of a single dict.

A notable unsupported data type is the `namedtuple`. The reason is that it behaves like both an ordered datatype (tuple) and a mapping datatype (dict). So given a `namedtuple` of the form:

```
`namedtuple("example_tuple", ["y", "x"])`
```

it is ambiguous whether to reverse the order of the elements when interpreting the value. Even worse is a tuple of the form:

```
`namedtuple("other_tuple", ["x", "y", "z"])`
```

where it is unclear if the tuple was intended to be unpacked into `x`, `y`, and `sample_weight` or passed through as a single element to ``x``. As a result the data processing code will simply raise a `ValueError` if it encounters a `namedtuple`. (Along with instructions to remedy the issue.)

Returns:

A ``History`` object. Its ``History.history`` attribute is a record of training loss values and metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

Raises:

`RuntimeError`: 1. If the model was never compiled or,
2. If ``model.fit`` is wrapped in ``tf.function``.

`ValueError`: In case of mismatch between the provided input data

and what the model expects or when the input data is empty.

`fit_generator(self, generator, steps_per_epoch=None, epochs=1, verbose=1, callbacks=None, validation_data=None, validation_steps=None, validation_freq=1, class_weight=None, max_queue_size=10, workers=1, use_multiprocessing=False, shuffle=True, initial_epoch=0)`
 Fits the model on data yielded batch-by-batch by a Python generator.

DEPRECATED:
 `Model.fit` now supports generators, so there is no longer any need to use this endpoint.

`get_layer(self, name=None, index=None)`
 Retrieves a layer based on either its name (unique) or index.

If `name` and `index` are both provided, `index` will take precedence. Indices are based on order of horizontal graph traversal (bottom-up).

Args:
 name: String, name of layer.
 index: Integer, index of layer.

Returns:
 A layer instance.

`get_weights(self)`
 Retrieves the weights of the model.

Returns:
 A flat list of Numpy arrays.

`load_weights(self, filepath, by_name=False, skip_mismatch=False, options=None)`
 Loads all layer weights, either from a TensorFlow or an HDF5 weight file.

If `by_name` is False weights are loaded based on the network's topology. This means the architecture should be the same as when the weights were saved. Note that layers that don't have weights are not taken into account in the topological ordering, so adding or removing layers is fine as long as they don't have weights.

If `by_name` is True, weights are loaded into layers only if they share the same name. This is useful for fine-tuning or transfer-learning models where some of the layers have changed.

Only topological loading (`by_name=False`) is supported when loading weights from the TensorFlow format. Note that topological loading differs slightly between TensorFlow and HDF5 formats for user-defined classes inheriting from `tf.keras.Model`: HDF5 loads based on a flattened list of weights, while the TensorFlow format loads based on the object-local names of attributes to which layers are assigned in the `Model`'s constructor.

Args:
 filepath: String, path to the weights file to load. For weight files in TensorFlow format, this is the file prefix (the same as was passed to `save_weights`). This can also be a path to a SavedModel saved from `model.save`.
 by_name: Boolean, whether to load weights by name or by topological order. Only topological loading is supported for weight files in

TensorFlow format.

skip_mismatch: Boolean, whether to skip loading of layers where there is a mismatch in the number of weights, or a mismatch in the shape of the weight (only valid when `by_name=True`).

options: Optional `tf.train.CheckpointOptions` object that specifies options for loading weights.

Returns:

When loading a weight file in TensorFlow format, returns the same status object as `tf.train.Checkpoint.restore`. When graph building, restore ops are run automatically as soon as the network is built (on first call for user-defined classes inheriting from `Model`, immediately if it is already built).

When loading weights in HDF5 format, returns `None`.

Raises:

ImportError: If `h5py` is not available and the weight file is in HDF5 format.

ValueError: If `skip_mismatch` is set to `True` when `by_name` is `False`.

`make_predict_function(self, force=False)`

Creates a function that executes one step of inference.

This method can be overridden to support custom inference logic. This method is called by `Model.predict` and `Model.predict_on_batch`.

Typically, this method directly controls `tf.function` and `tf.distribute.Strategy` settings, and delegates the actual evaluation logic to `Model.predict_step`.

This function is cached the first time `Model.predict` or `Model.predict_on_batch` is called. The cache is cleared whenever `Model.compile` is called. You can skip the cache and generate again the function with `force=True`.

Args:

force: Whether to regenerate the predict function and skip the cached function if available.

Returns:

Function. The function created by this method should accept a `tf.data.Iterator`, and return the outputs of the `Model`.

`make_test_function(self, force=False)`

Creates a function that executes one step of evaluation.

This method can be overridden to support custom evaluation logic. This method is called by `Model.evaluate` and `Model.test_on_batch`.

Typically, this method directly controls `tf.function` and `tf.distribute.Strategy` settings, and delegates the actual evaluation logic to `Model.test_step`.

This function is cached the first time `Model.evaluate` or `Model.test_on_batch` is called. The cache is cleared whenever `Model.compile` is called. You can skip the cache and generate again the function with `force=True`.

Args:

`force`: Whether to regenerate the test function and skip the cached function if available.

Returns:

Function. The function created by this method should accept a ``tf.data.Iterator``, and return a ``dict`` containing values that will be passed to ``tf.keras.Callbacks.on_test_batch_end``.

`make_train_function(self, force=False)`

Creates a function that executes one step of training.

This method can be overridden to support custom training logic. This method is called by ``Model.fit`` and ``Model.train_on_batch``.

Typically, this method directly controls ``tf.function`` and ``tf.distribute.Strategy`` settings, and delegates the actual training logic to ``Model.train_step``.

This function is cached the first time ``Model.fit`` or ``Model.train_on_batch`` is called. The cache is cleared whenever ``Model.compile`` is called. You can skip the cache and generate again the function with ``force=True``.

Args:

`force`: Whether to regenerate the train function and skip the cached function if available.

Returns:

Function. The function created by this method should accept a ``tf.data.Iterator``, and return a ``dict`` containing values that will be passed to ``tf.keras.Callbacks.on_train_batch_end``, such as ``{'loss': 0.2, 'accuracy': 0.7}``.

`predict(self, x, batch_size=None, verbose='auto', steps=None, callbacks=None, max_queue_size=10, workers=1, use_multiprocessing=False)`

Generates output predictions for the input samples.

Computation is done in batches. This method is designed for batch processing of large numbers of inputs. It is not intended for use inside of loops that iterate over your data and process small numbers of inputs at a time.

For small numbers of inputs that fit in one batch, directly use ``__call__()`` for faster execution, e.g., ``model(x)``, or ``model(x, training=False)`` if you have layers such as ``tf.keras.layers.BatchNormalization`` that behave differently during inference. You may pair the individual model call with a ``tf.function`` for additional performance inside your inner loop. If you need access to numpy array values instead of tensors after your model call, you can use ``tensor.numpy()`` to get the numpy array value of an eager tensor.

Also, note the fact that test loss is not affected by regularization layers like noise and dropout.

Note: See [this FAQ entry](https://keras.io/getting_started/faq/#whats-the-difference-between-model-methods-predict-and-call) for more details about the difference between ``Model`` methods

``predict()`` and ``__call__()``.

Args:

`x`: Input samples. It could be:

- A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
- A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).
- A ``tf.data`` dataset.
- A generator or ``keras.utils.Sequence`` instance.

A more detailed description of unpacking behavior for iterator types (Dataset, generator, Sequence) is given in the ``Unpacking behavior for iterator-like inputs`` section of ``Model.fit``.

`batch_size`: Integer or ``None``.

Number of samples per batch.

If unspecified, ``batch_size`` will default to 32.

Do not specify the ``batch_size`` if your data is in the form of dataset, generators, or ``keras.utils.Sequence`` instances (since they generate batches).

`verbose`: ``"auto"`, 0, 1, or 2`. Verbosity mode.

0 = silent, 1 = progress bar, 2 = single line.

``"auto"`, defaults to 1 for most cases, and to 2 when used with `ParameterServerStrategy`. Note that the progress bar is not particularly useful when logged to a file, so `verbose=2` is recommended when not running interactively (e.g. in a production environment).`

`steps`: Total number of steps (batches of samples) before declaring the prediction round finished.

Ignored with the default value of ``None``. If `x` is a ``tf.data`` dataset and ``steps`` is ``None``, ``predict()`` will run until the input dataset is exhausted.

`callbacks`: List of ``keras.callbacks.Callback`` instances.

List of callbacks to apply during prediction.

See [\[callbacks\]\(/api_docs/python/tf/keras/callbacks\)](#).

`max_queue_size`: Integer. Used for generator or

``keras.utils.Sequence`` input only. Maximum size for the generator queue. If unspecified, ``max_queue_size`` will default to 10.

`workers`: Integer. Used for generator or ``keras.utils.Sequence`` input only. Maximum number of processes to spin up when using process-based threading. If unspecified, ``workers`` will default to 1.

`use_multiprocessing`: Boolean. Used for generator or

``keras.utils.Sequence`` input only. If ``True``, use process-based threading. If unspecified, ``use_multiprocessing`` will default to ``False``. Note that because this implementation relies on multiprocessing, you should not pass non-picklable arguments to the generator as they can't be passed easily to children processes.

See the discussion of ``Unpacking behavior for iterator-like inputs`` for ``Model.fit``. Note that `Model.predict` uses the same interpretation rules as ``Model.fit`` and ``Model.evaluate``, so inputs must be unambiguous for all three methods.

Returns:

Numpy array(s) of predictions.

Raises:

`RuntimeError`: If ``model.predict`` is wrapped in a ``tf.function``.

ValueError: In case of mismatch between the provided input data and the model's expectations, or in case a stateful model receives a number of samples that is not a multiple of the batch size.

`predict_generator(self, generator, steps=None, callbacks=None, max_queue_size=10, workers=1, use_multiprocessing=False, verbose=0)`

Generates predictions for the input samples from a data generator.

DEPRECATED:

``Model.predict`` now supports generators, so there is no longer any need to use this endpoint.

`predict_on_batch(self, x)`

Returns predictions for a single batch of samples.

Args:

x: Input data. It could be:

- A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
- A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).

Returns:

Numpy array(s) of predictions.

Raises:

RuntimeError: If ``model.predict_on_batch`` is wrapped in a ``tf.function``.

`predict_step(self, data)`

The logic for one inference step.

This method can be overridden to support custom inference logic.

This method is called by ``Model.make_predict_function``.

This method should contain the mathematical logic for one step of inference. This typically includes the forward pass.

Configuration details for *how* this logic is run (e.g. ``tf.function`` and ``tf.distribute.Strategy`` settings), should be left to ``Model.make_predict_function``, which can also be overridden.

Args:

data: A nested structure of ``Tensor`s`.

Returns:

The result of one inference step, typically the output of calling the ``Model`` on data.

`reset_metrics(self)`

Resets the state of all the metrics in the model.

Examples:

```
>>> inputs = tf.keras.layers.Input(shape=(3,))
>>> outputs = tf.keras.layers.Dense(2)(inputs)
>>> model = tf.keras.models.Model(inputs=inputs, outputs=outputs)
>>> model.compile(optimizer="Adam", loss="mse", metrics=["mae"])
```

```

>>> x = np.random.random((2, 3))
>>> y = np.random.randint(0, 2, (2, 2))
>>> _ = model.fit(x, y, verbose=0)
>>> assert all(float(m.result()) for m in model.metrics)

>>> model.reset_metrics()
>>> assert all(float(m.result()) == 0 for m in model.metrics)

reset_states(self)

save(self, filepath, overwrite=True, include_optimizer=True, save_format=None,
e, signatures=None, options=None, save_traces=True)
    Saves the model to Tensorflow SavedModel or a single HDF5 file.

    Please see `tf.keras.models.save_model` or the
    [Serialization and Saving guide](
    https://keras.io/guides/serialization_and_saving/)
    for details.

    Args:
        filepath: String, PathLike, path to SavedModel or H5 file to save
            the model.
        overwrite: Whether to silently overwrite any existing file at the
            target location, or provide the user with a manual prompt.
        include_optimizer: If True, save optimizer's state together.
        save_format: Either `tf` or `h5`, indicating whether to save the
            model to Tensorflow SavedModel or HDF5. Defaults to 'tf' in TF
            2.X, and 'h5' in TF 1.X.
        signatures: Signatures to save with the SavedModel. Applicable to
            the 'tf' format only. Please see the `signatures` argument in
            `tf.saved_model.save` for details.
        options: (only applies to SavedModel format)
            `tf.saved_model.SaveOptions` object that specifies options for
            saving to SavedModel.
        save_traces: (only applies to SavedModel format) When enabled, the
            SavedModel will store the function traces for each layer. This
            can be disabled, so that only the configs of each layer are
            stored. Defaults to `True`. Disabling this will decrease
            serialization time and reduce file size, but it requires that
            all custom layers/models implement a `get_config()` method.

    Example:

    ```python
 from keras.models import load_model

 model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
 del model # deletes the existing model

 # returns a compiled model
 # identical to the previous one
 model = load_model('my_model.h5')
    ```

    save_spec(self, dynamic_batch=True)
        Returns the `tf.TensorSpec` of call inputs as a tuple `(args, kwargs)`.

        This value is automatically defined after calling the model for the
        first time. Afterwards, you can use it when exporting the model for
        serving:

```



```

'''python
model = tf.keras.Model(...)

@tf.function
def serve(*args, **kwargs):
    outputs = model(*args, **kwargs)
    # Apply postprocessing steps, or add additional outputs.
    ...
    return outputs

# arg_specs is `[tf.TensorSpec(...), ...]`. kwarg_specs, in this
# example, is an empty dict since functional models do not use keyword
# arguments.
arg_specs, kwarg_specs = model.save_spec()

model.save(path, signatures={
    'serving_default': serve.get_concrete_function(*arg_specs,
                                                    **kwarg_specs)
})
'''

```

Args:

`dynamic_batch`: Whether to set the batch sizes of all the returned `tf.TensorSpec` to `None`. (Note that when defining functional or Sequential models with `tf.keras.Input(..., batch_size=X)`, the batch size will always be preserved). Defaults to `True`.

Returns:

If the model inputs are defined, returns a tuple `(args, kwargs)`. All elements in `args` and `kwargs` are `tf.TensorSpec`.
 If the model inputs are not defined, returns `None`.
 The model inputs are automatically set when calling the model, `model.fit`, `model.evaluate` or `model.predict`.

`save_weights(self, filepath, overwrite=True, save_format=None, options=None)`
 Saves all layer weights.

Either saves in HDF5 or in TensorFlow format based on the `save_format` argument.

When saving in HDF5 format, the weight file has:

- `layer_names` (attribute), a list of strings (ordered names of model layers).
- For every layer, a `group` named `layer.name`
 - For every such layer group, a group attribute `weight_names`, a list of strings (ordered names of weights tensor of the layer).
 - For every weight in the layer, a dataset storing the weight value, named after the weight tensor.

When saving in TensorFlow format, all objects referenced by the network are saved in the same format as `tf.train.Checkpoint`, including any `Layer` instances or `Optimizer` instances assigned to object attributes. For networks constructed from inputs and outputs using `tf.keras.Model(inputs, outputs)`, `Layer` instances used by the network are tracked/saved automatically. For user-defined classes which inherit from `tf.keras.Model`, `Layer` instances must be assigned to object attributes, typically in the constructor. See the documentation of `tf.train.Checkpoint` and `tf.keras.Model` for details.

While the formats are the same, do not mix ``save_weights`` and ``tf.train.Checkpoint``. Checkpoints saved by ``Model.save_weights`` should be loaded using ``Model.load_weights``. Checkpoints saved using ``tf.train.Checkpoint.save`` should be restored using the corresponding ``tf.train.Checkpoint.restore``. Prefer ``tf.train.Checkpoint`` over ``save_weights`` for training checkpoints.

The TensorFlow format matches objects and variables by starting at a root object, ``self`` for ``save_weights``, and greedily matching attribute names. For ``Model.save`` this is the ``Model``, and for ``Checkpoint.save`` this is the ``Checkpoint`` even if the ``Checkpoint`` has a model attached. This means saving a ``tf.keras.Model`` using ``save_weights`` and loading into a ``tf.train.Checkpoint`` with a ``Model`` attached (or vice versa) will not match the ``Model``'s variables. See the [guide to training checkpoints](<https://www.tensorflow.org/guide/checkpoint>) for details on the TensorFlow format.

Args:

`filepath`: String or PathLike, path to the file to save the weights to. When saving in TensorFlow format, this is the prefix used for checkpoint files (multiple files are generated). Note that the `'.h5'` suffix causes weights to be saved in HDF5 format.

`overwrite`: Whether to silently overwrite any existing file at the target location, or provide the user with a manual prompt.

`save_format`: Either `'tf'` or `'h5'`. A ``filepath`` ending in `'.h5'` or `'.keras'` will default to HDF5 if ``save_format`` is ``None``. Otherwise ``None`` defaults to `'tf'`.

`options`: Optional ``tf.train.CheckpointOptions`` object that specifies options for saving weights.

Raises:

`ImportError`: If ``h5py`` is not available when attempting to save in HDF5 format.

`summary(self, line_length=None, positions=None, print_fn=None, expand_nested=False, show_trainable=False, layer_range=None)`
Prints a string summary of the network.

Args:

`line_length`: Total length of printed lines (e.g. set this to adapt the display to different terminal window sizes).

`positions`: Relative or absolute positions of log elements in each line. If not provided, defaults to ``[.33, .55, .67, 1.]``.

`print_fn`: Print function to use. Defaults to ``print``. It will be called on each line of the summary. You can set it to a custom function in order to capture the string summary.

`expand_nested`: Whether to expand the nested models. If not provided, defaults to ``False``.

`show_trainable`: Whether to show if a layer is trainable. If not provided, defaults to ``False``.

`layer_range`: a list or tuple of 2 strings, which is the starting layer name and ending layer name (both inclusive) indicating the range of layers to be printed in summary. It also accepts regex patterns instead of exact name. In such case, start predicate will be the first element it matches to ``layer_range[0]`` and the end predicate will be

the last element it matches to ``layer_range[1]``.
By default ``None`` which considers all layers of model.

Raises:

`ValueError`: if ``summary()`` is called before the model is built.

`test_on_batch(self, x, y=None, sample_weight=None, reset_metrics=True, return_dict=False)`

Test the model on a single batch of samples.

Args:

`x`: Input data. It could be:

- A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
- A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).
- A dict mapping input names to the corresponding array/tensors, if the model has named inputs.

`y`: Target data. Like the input data ``x``, it could be either Numpy array(s) or TensorFlow tensor(s). It should be consistent with ``x`` (you cannot have Numpy inputs and tensor targets, or inversely).

`sample_weight`: Optional array of the same length as `x`, containing weights to apply to the model's loss for each sample. In the case of temporal data, you can pass a 2D array with shape (samples, sequence_length), to apply a different weight to every timestep of every sample.

`reset_metrics`: If ``True``, the metrics returned will be only for this batch. If ``False``, the metrics will be statefully accumulated across batches.

`return_dict`: If ``True``, loss and metric results are returned as a dict, with each key being the name of the metric. If ``False``, they are returned as a list.

Returns:

Scalar test loss (if the model has a single output and no metrics) or list of scalars (if the model has multiple outputs and/or metrics). The attribute ``model.metrics_names`` will give you the display labels for the scalar outputs.

Raises:

`RuntimeError`: If ``model.test_on_batch`` is wrapped in a ``tf.function``.

`test_step(self, data)`

The logic for one evaluation step.

This method can be overridden to support custom evaluation logic.
This method is called by ``Model.make_test_function``.

This function should contain the mathematical logic for one step of evaluation.

This typically includes the forward pass, loss calculation, and metrics updates.

Configuration details for *how* this logic is run (e.g. ``tf.function`` and ``tf.distribute.Strategy`` settings), should be left to ``Model.make_test_function``, which can also be overridden.

Args:

`data`: A nested structure of ``Tensor`s`.

Returns:

A `dict` containing values that will be passed to `tf.keras.callbacks.CallbackList.on_train_batch_end`. Typically, the values of the `Model`'s metrics are returned.

to_json(self, **kwargs)

Returns a JSON string containing the network configuration.

To load a network from a JSON save file, use

`keras.models.model_from_json(json_string, custom_objects={})`.

Args:

****kwargs:** Additional keyword arguments to be passed to `*`json.dumps()``.

Returns:

A JSON string.

to_yaml(self, **kwargs)

Returns a yaml string containing the network configuration.

Note: Since TF 2.6, this method is no longer supported and will raise a `RuntimeError`.

To load a network from a yaml save file, use

`keras.models.model_from_yaml(yaml_string, custom_objects={})`.

`custom_objects` should be a dictionary mapping the names of custom losses / layers / etc to the corresponding functions / classes.

Args:

****kwargs:** Additional keyword arguments to be passed to ``yaml.dump()``.

Returns:

A YAML string.

Raises:

`RuntimeError`: announces that the method poses a security risk

train_on_batch(self, x, y=None, sample_weight=None, class_weight=None, reset_metrics=True, return_dict=False)

Runs a single gradient update on a single batch of data.

Args:

x: Input data. It could be:

- A Numpy array (or array-like), or a list of arrays (in case the model has multiple inputs).
- A TensorFlow tensor, or a list of tensors (in case the model has multiple inputs).
- A dict mapping input names to the corresponding array/tensors, if the model has named inputs.

y: Target data. Like the input data `x`, it could be either Numpy array(s) or TensorFlow tensor(s).

sample_weight: Optional array of the same length as x, containing weights to apply to the model's loss for each sample. In the case of temporal data, you can pass a 2D array with shape (samples, sequence_length), to apply a different weight to every timestep of

every sample.

`class_weight`: Optional dictionary mapping class indices (integers) to a weight (float) to apply to the model's loss for the samples from this class during training. This can be useful to tell the model to "pay more attention" to samples from an under-represented class.

`reset_metrics`: If ``True``, the metrics returned will be only for this batch. If ``False``, the metrics will be statefully accumulated across batches.

`return_dict`: If ``True``, loss and metric results are returned as a dict, with each key being the name of the metric. If ``False``, they are returned as a list.

Returns:

Scalar training loss
(if the model has a single output and no metrics)
or list of scalars (if the model has multiple outputs and/or metrics). The attribute ``model.metrics_names`` will give you the display labels for the scalar outputs.

Raises:

`RuntimeError`: If ``model.train_on_batch`` is wrapped in a ``tf.function``.

`train_step(self, data)`

The logic for one training step.

This method can be overridden to support custom training logic. For concrete examples of how to override this method see [Customizing what happens in fit](https://www.tensorflow.org/guide/keras/customizing_what_happens_in_fit). This method is called by ``Model.make_train_function``.

This method should contain the mathematical logic for one step of training. This typically includes the forward pass, loss calculation, backpropagation, and metric updates.

Configuration details for *how* this logic is run (e.g. ``tf.function`` and ``tf.distribute.Strategy`` settings), should be left to ``Model.make_train_function``, which can also be overridden.

Args:

`data`: A nested structure of ``Tensor``s.

Returns:

A ``dict`` containing values that will be passed to ``tf.keras.callbacks.CallbackList.on_train_batch_end``. Typically, the values of the ``Model``'s metrics are returned. Example:
``{'loss': 0.2, 'accuracy': 0.7}``.

Static methods inherited from `keras.engine.training.Model`:

`__new__(cls, *args, **kwargs)`

Create and return a new object. See `help(type)` for accurate signature.

Readonly properties inherited from `keras.engine.training.Model`:

`distribute_strategy`

The ``tf.distribute.Strategy`` this model was created under.

metrics

Returns the model's metrics added using ``compile()``, ``add_metric()`` APIs.

Note: Metrics passed to ``compile()`` are available only after a ``keras.Model`` has been trained/evaluated on actual data.

Examples:

```
>>> inputs = tf.keras.layers.Input(shape=(3,))
>>> outputs = tf.keras.layers.Dense(2)(inputs)
>>> model = tf.keras.models.Model(inputs=inputs, outputs=outputs)
>>> model.compile(optimizer="Adam", loss="mse", metrics=["mae"])
>>> [m.name for m in model.metrics]
[]

>>> x = np.random.random((2, 3))
>>> y = np.random.randint(0, 2, (2, 2))
>>> model.fit(x, y)
>>> [m.name for m in model.metrics]
['loss', 'mae']

>>> inputs = tf.keras.layers.Input(shape=(3,))
>>> d = tf.keras.layers.Dense(2, name='out')
>>> output_1 = d(inputs)
>>> output_2 = d(inputs)
>>> model = tf.keras.models.Model(
...     inputs=inputs, outputs=[output_1, output_2])
>>> model.add_metric(
...     tf.reduce_sum(output_2), name='mean', aggregation='mean')
>>> model.compile(optimizer="Adam", loss="mse", metrics=["mae", "acc"])
>>> model.fit(x, (y, y))
>>> [m.name for m in model.metrics]
['loss', 'out_loss', 'out_1_loss', 'out_mae', 'out_acc', 'out_1_mae',
'out_1_acc', 'mean']
```

metrics_names

Returns the model's display labels for all outputs.

Note: ``metrics_names`` are available only after a ``keras.Model`` has been trained/evaluated on actual data.

Examples:

```
>>> inputs = tf.keras.layers.Input(shape=(3,))
>>> outputs = tf.keras.layers.Dense(2)(inputs)
>>> model = tf.keras.models.Model(inputs=inputs, outputs=outputs)
>>> model.compile(optimizer="Adam", loss="mse", metrics=["mae"])
>>> model.metrics_names
[]

>>> x = np.random.random((2, 3))
>>> y = np.random.randint(0, 2, (2, 2))
>>> model.fit(x, y)
>>> model.metrics_names
['loss', 'mae']

>>> inputs = tf.keras.layers.Input(shape=(3,))
>>> d = tf.keras.layers.Dense(2, name='out')
>>> output_1 = d(inputs)
```

```
>>> output_2 = d(inputs)
>>> model = tf.keras.models.Model(
...     inputs=inputs, outputs=[output_1, output_2])
>>> model.compile(optimizer="Adam", loss="mse", metrics=["mae", "acc"])
>>> model.fit(x, (y, y))
>>> model.metrics_names
['loss', 'out_loss', 'out_1_loss', 'out_mae', 'out_acc', 'out_1_mae',
'out_1_acc']
```

non_trainable_weights

List of all non-trainable weights tracked by this layer.

Non-trainable weights are **not** updated during training. They are expected to be updated manually in `call()`.

Returns:

A list of non-trainable variables.

state_updates

Deprecated, do NOT use!

Returns the `updates` from all layers that are stateful.

This is useful for separating training updates and state updates, e.g. when we need to update a layer's internal state during prediction.

Returns:

A list of update ops.

trainable_weights

List of all trainable weights tracked by this layer.

Trainable weights are updated via gradient descent during training.

Returns:

A list of trainable variables.

weights

Returns the list of all layer variables/weights.

Note: This will not track the weights of nested `tf.Modules` that are not themselves Keras layers.

Returns:

A list of variables.

Data descriptors inherited from `keras.engine.training.Model`:

run_eagerly

Settable attribute indicating whether the model should run eagerly.

Running eagerly means that your model will be run step by step, like Python code. Your model might run slower, but it should become easier for you to debug it by stepping into individual layer calls.

By default, we will attempt to compile your model to a static graph to deliver the best execution performance.

Returns:

Boolean, whether the model should run eagerly.

Methods inherited from keras.engine.base_layer.Layer:

`__delattr__(self, name)`
Implement `delattr(self, name)`.

`__getstate__(self)`

`__setstate__(self, state)`

`add_loss(self, losses, **kwargs)`
Add loss tensor(s), potentially dependent on layer inputs.

Some losses (for instance, activity regularization losses) may be dependent on the inputs passed when calling a layer. Hence, when reusing the same layer on different inputs ``a`` and ``b``, some entries in ``layer.losses`` may be dependent on ``a`` and some on ``b``. This method automatically keeps track of dependencies.

This method can be used inside a subclassed layer or model's ``call`` function, in which case ``losses`` should be a Tensor or list of Tensors.

Example:

```
```python
class MyLayer(tf.keras.layers.Layer):
 def call(self, inputs):
 self.add_loss(tf.abs(tf.reduce_mean(inputs)))
 return inputs
```
```

This method can also be called directly on a Functional Model during construction. In this case, any loss Tensors passed to this Model must be symbolic and be able to be traced back to the model's ``Input``s. These losses become part of the model's topology and are tracked in ``get_config``.

Example:

```
```python
inputs = tf.keras.Input(shape=(10,))
x = tf.keras.layers.Dense(10)(inputs)
outputs = tf.keras.layers.Dense(1)(x)
model = tf.keras.Model(inputs, outputs)
Activity regularization.
model.add_loss(tf.abs(tf.reduce_mean(x)))
```
```

If this is not the case for your loss (if, for example, your loss references a ``Variable`` of one of the model's layers), you can wrap your loss in a zero-argument lambda. These losses are not tracked as part of the model's topology since they can't be serialized.

Example:

```
```python
inputs = tf.keras.Input(shape=(10,))
```



```

d = tf.keras.layers.Dense(10)
x = d(inputs)
outputs = tf.keras.layers.Dense(1)(x)
model = tf.keras.Model(inputs, outputs)
Weight regularization.
model.add_loss(lambda: tf.reduce_mean(d.kernel))
'''

```

Args:

losses: Loss tensor, or list/tuple of tensors. Rather than tensors, losses may also be zero-argument callables which create a loss tensor.

\*\*kwargs: Used for backwards compatibility only.

```

add_metric(self, value, name=None, **kwargs)
 Adds metric tensor to the layer.

```

This method can be used inside the `call()` method of a subclassed layer or model.

```

'''python
class MyMetricLayer(tf.keras.layers.Layer):
 def __init__(self):
 super(MyMetricLayer, self).__init__(name='my_metric_layer')
 self.mean = tf.keras.metrics.Mean(name='metric_1')

 def call(self, inputs):
 self.add_metric(self.mean(inputs))
 self.add_metric(tf.reduce_sum(inputs), name='metric_2')
 return inputs
'''

```

This method can also be called directly on a Functional Model during construction. In this case, any tensor passed to this Model must be symbolic and be able to be traced back to the model's `Input`'s. These metrics become part of the model's topology and are tracked when you save the model via `save()`.

```

'''python
inputs = tf.keras.Input(shape=(10,))
x = tf.keras.layers.Dense(10)(inputs)
outputs = tf.keras.layers.Dense(1)(x)
model = tf.keras.Model(inputs, outputs)
model.add_metric(math_ops.reduce_sum(x), name='metric_1')
'''

```

Note: Calling `add_metric()` with the result of a metric object on a Functional Model, as shown in the example below, is not supported. This is because we cannot trace the metric result tensor back to the model's inputs.

```

'''python
inputs = tf.keras.Input(shape=(10,))
x = tf.keras.layers.Dense(10)(inputs)
outputs = tf.keras.layers.Dense(1)(x)
model = tf.keras.Model(inputs, outputs)
model.add_metric(tf.keras.metrics.Mean()(x), name='metric_1')
'''

```

Args:

```

 value: Metric tensor.
 name: String metric name.
 **kwargs: Additional keyword arguments for backward compatibility.
 Accepted values:
 `aggregation` - When the `value` tensor provided is not the result
 of calling a `keras.Metric` instance, it will be aggregated by
 default using a `keras.Metric.Mean`.

```

`add_update(self, updates)`  
 Add update op(s), potentially dependent on layer inputs.

Weight updates (for instance, the updates of the moving mean and variance in a BatchNormalization layer) may be dependent on the inputs passed when calling a layer. Hence, when reusing the same layer on different inputs `a` and `b`, some entries in `layer.updates` may be dependent on `a` and some on `b`. This method automatically keeps track of dependencies.

This call is ignored when eager execution is enabled (in that case, variable updates are run on the fly and thus do not need to be tracked for later execution).

Args:

- `updates`: Update op, or list/tuple of update ops, or zero-arg callable that returns an update op. A zero-arg callable should be passed in order to disable running the updates by setting `trainable=False` on this Layer, when executing in Eager mode.

```

 add_variable(self, *args, **kwargs)
 Deprecated, do NOT use! Alias for `add_weight`.

```

```

 add_weight(self, name=None, shape=None, dtype=None, initializer=None, regular
izer=None, trainable=None, constraint=None, use_resource=None, synchronization=<V
ariableSynchronization.AUTO: 0>, aggregation=<VariableAggregationV2.NONE: 0>, **k
wargs)
 Adds a new variable to the layer.

```

Args:

- `name`: Variable name.
- `shape`: Variable shape. Defaults to scalar if unspecified.
- `dtype`: The type of the variable. Defaults to `self.dtype`.
- `initializer`: Initializer instance (callable).
- `regularizer`: Regularizer instance (callable).
- `trainable`: Boolean, whether the variable should be part of the layer's "trainable\_variables" (e.g. variables, biases) or "non\_trainable\_variables" (e.g. BatchNorm mean and variance). Note that `trainable` cannot be `True` if `synchronization` is set to `ON\_READ`.
- `constraint`: Constraint instance (callable).
- `use_resource`: Whether to use a `ResourceVariable` or not. See [this guide]([https://www.tensorflow.org/guide/migrate/tf1\\_vs\\_tf2#resourcevariables\\_instead\\_of\\_referencevariables](https://www.tensorflow.org/guide/migrate/tf1_vs_tf2#resourcevariables_instead_of_referencevariables)) for more information.
- `synchronization`: Indicates when a distributed a variable will be aggregated. Accepted values are constants defined in the class `tf.VariableSynchronization`. By default the synchronization is set to `AUTO` and the current `DistributionStrategy` chooses when to synchronize. If `synchronization` is set to `ON\_READ`, `trainable` must not be set to `True`.

aggregation: Indicates how a distributed variable will be aggregated. Accepted values are constants defined in the class ``tf.VariableAggregation``.  
 \*\*kwargs: Additional keyword arguments. Accepted values are ``getter``, ``collections``, ``experimental_autocast`` and ``caching_device``.

Returns:  
 The variable created.

Raises:  
 ValueError: When giving unsupported dtype and no initializer or when trainable has been set to True with synchronization set as ``ON_READ``.

`compute_output_signature(self, input_signature)`  
 Compute the output tensor signature of the layer based on the inputs.

Unlike a `TensorShape` object, a `TensorSpec` object contains both shape and dtype information for a tensor. This method allows layers to provide output dtype information if it is different from the input dtype. For any layer that doesn't implement this function, the framework will fall back to use ``compute_output_shape``, and will assume that the output dtype matches the input dtype.

Args:  
 input\_signature: Single `TensorSpec` or nested structure of `TensorSpec` objects, describing a candidate input for the layer.

Returns:  
 Single `TensorSpec` or nested structure of `TensorSpec` objects, describing how the layer would transform the provided input.

Raises:  
 TypeError: If `input_signature` contains a non-`TensorSpec` object.

`count_params(self)`  
 Count the total number of scalars composing the weights.

Returns:  
 An integer count.

Raises:  
 ValueError: if the layer isn't yet built (in which case its weights aren't yet defined).

`finalize_state(self)`  
 Finalizes the layers state after updating layer weights.

This function can be subclassed in a layer and will be called after updating a layer weights. It can be overridden to finalize any additional layer state after a weight update.

This function will be called after weights of a layer have been restored from a loaded model.

`get_input_at(self, node_index)`  
 Retrieves the input tensor(s) of a layer at a given node.

Args:  
 node\_index: Integer, index of the node

from which to retrieve the attribute.  
 E.g. `node\_index=0` will correspond to the  
 first input node of the layer.

Returns:

A tensor (or list of tensors if the layer has multiple inputs).

Raises:

RuntimeError: If called in Eager mode.

`get_input_mask_at(self, node_index)`

Retrieves the input mask tensor(s) of a layer at a given node.

Args:

node\_index: Integer, index of the node  
 from which to retrieve the attribute.  
 E.g. `node\_index=0` will correspond to the  
 first time the layer was called.

Returns:

A mask tensor  
 (or list of tensors if the layer has multiple inputs).

`get_input_shape_at(self, node_index)`

Retrieves the input shape(s) of a layer at a given node.

Args:

node\_index: Integer, index of the node  
 from which to retrieve the attribute.  
 E.g. `node\_index=0` will correspond to the  
 first time the layer was called.

Returns:

A shape tuple  
 (or list of shape tuples if the layer has multiple inputs).

Raises:

RuntimeError: If called in Eager mode.

`get_output_at(self, node_index)`

Retrieves the output tensor(s) of a layer at a given node.

Args:

node\_index: Integer, index of the node  
 from which to retrieve the attribute.  
 E.g. `node\_index=0` will correspond to the  
 first output node of the layer.

Returns:

A tensor (or list of tensors if the layer has multiple outputs).

Raises:

RuntimeError: If called in Eager mode.

`get_output_mask_at(self, node_index)`

Retrieves the output mask tensor(s) of a layer at a given node.

Args:

node\_index: Integer, index of the node  
 from which to retrieve the attribute.

E.g. `node\_index=0` will correspond to the first time the layer was called.

Returns:

A mask tensor  
(or list of tensors if the layer has multiple outputs).

`get_output_shape_at(self, node_index)`

Retrieves the output shape(s) of a layer at a given node.

Args:

`node_index`: Integer, index of the node from which to retrieve the attribute.  
E.g. `node\_index=0` will correspond to the first time the layer was called.

Returns:

A shape tuple  
(or list of shape tuples if the layer has multiple outputs).

Raises:

`RuntimeError`: If called in Eager mode.

`set_weights(self, weights)`

Sets the weights of the layer, from NumPy arrays.

The weights of a layer represent the state of the layer. This function sets the weight values from numpy arrays. The weight values should be passed in the order they are created by the layer. Note that the layer's weights must be instantiated before calling this function, by calling the layer.

For example, a `Dense` layer returns a list of two values: the kernel matrix and the bias vector. These can be used to set the weights of another `Dense` layer:

```
>>> layer_a = tf.keras.layers.Dense(1,
... kernel_initializer=tf.constant_initializer(1.))
>>> a_out = layer_a(tf.convert_to_tensor([[1., 2., 3.])))
>>> layer_a.get_weights()
[array([[1.],
 [1.],
 [1.]], dtype=float32), array([0.], dtype=float32)]
>>> layer_b = tf.keras.layers.Dense(1,
... kernel_initializer=tf.constant_initializer(2.))
>>> b_out = layer_b(tf.convert_to_tensor([[10., 20., 30.])))
>>> layer_b.get_weights()
[array([[2.],
 [2.],
 [2.]], dtype=float32), array([0.], dtype=float32)]
>>> layer_b.set_weights(layer_a.get_weights())
>>> layer_b.get_weights()
[array([[1.],
 [1.],
 [1.]], dtype=float32), array([0.], dtype=float32)]
```

Args:

`weights`: a list of NumPy arrays. The number of arrays and their shape must match number of the dimensions of the weights

of the layer (i.e. it should match the output of ``get_weights``).

Raises:

`ValueError`: If the provided weights list does not match the layer's specifications.

-----  
Readonly properties inherited from `keras.engine.base_layer.Layer`:

`compute_dtype`

The dtype of the layer's computations.

This is equivalent to ``Layer.dtype_policy.compute_dtype``. Unless mixed precision is used, this is the same as ``Layer.dtype``, the dtype of the weights.

Layers automatically cast their inputs to the compute dtype, which causes computations and the output to be in the compute dtype as well. This is done by the base Layer class in ``Layer.__call__``, so you do not have to insert these casts if implementing your own layer.

Layers often perform certain internal computations in higher precision when ``compute_dtype`` is `float16` or `bfloat16` for numeric stability. The output will still typically be `float16` or `bfloat16` in such cases.

Returns:

The layer's compute dtype.

`dtype`

The dtype of the layer weights.

This is equivalent to ``Layer.dtype_policy.variable_dtype``. Unless mixed precision is used, this is the same as ``Layer.compute_dtype``, the dtype of the layer's computations.

`dtype_policy`

The dtype policy associated with this layer.

This is an instance of a ``tf.keras.mixed_precision.Policy``.

`dynamic`

Whether the layer is dynamic (eager-only); set in the constructor.

`inbound_nodes`

Return Functional API nodes upstream of this layer.

`input_mask`

Retrieves the input mask tensor(s) of a layer.

Only applicable if the layer has exactly one inbound node, i.e. if it is connected to one incoming layer.

Returns:

Input mask tensor (potentially `None`) or list of input mask tensors.

Raises:

`AttributeError`: if the layer is connected to more than one incoming layers.

**losses**

List of losses added using the ``add_loss()`` API.

Variable regularization tensors are created when this property is accessed, so it is eager safe: accessing ``losses`` under a ``tf.GradientTape`` will propagate gradients back to the corresponding variables.

Examples:

```
>>> class MyLayer(tf.keras.layers.Layer):
... def call(self, inputs):
... self.add_loss(tf.abs(tf.reduce_mean(inputs)))
... return inputs
>>> l = MyLayer()
>>> l(np.ones((10, 1)))
>>> l.losses
[1.0]

>>> inputs = tf.keras.Input(shape=(10,))
>>> x = tf.keras.layers.Dense(10)(inputs)
>>> outputs = tf.keras.layers.Dense(1)(x)
>>> model = tf.keras.Model(inputs, outputs)
>>> # Activity regularization.
>>> len(model.losses)
0
>>> model.add_loss(tf.abs(tf.reduce_mean(x)))
>>> len(model.losses)
1

>>> inputs = tf.keras.Input(shape=(10,))
>>> d = tf.keras.layers.Dense(10, kernel_initializer='ones')
>>> x = d(inputs)
>>> outputs = tf.keras.layers.Dense(1)(x)
>>> model = tf.keras.Model(inputs, outputs)
>>> # Weight regularization.
>>> model.add_loss(lambda: tf.reduce_mean(d.kernel))
>>> model.losses
[<tf.Tensor: shape=(), dtype=float32, numpy=1.0>]
```

Returns:

A list of tensors.

**name**

Name of the layer (string), set in the constructor.

**non\_trainable\_variables**

Sequence of non-trainable variables owned by this module and its submodules.

es.

Note: this method uses reflection to find variables on the current instance and submodules. For performance reasons you may wish to cache the result of calling this method if you don't expect the return value to change.

Returns:

A sequence of variables for the current module (sorted by attribute name) followed by variables from all submodules recursively (breadth first).

```

|
| outbound_nodes
| Return Functional API nodes downstream of this layer.
|
| output_mask
| Retrieves the output mask tensor(s) of a layer.
|
| Only applicable if the layer has exactly one inbound node,
| i.e. if it is connected to one incoming layer.
|
| Returns:
| Output mask tensor (potentially None) or list of output
| mask tensors.
|
| Raises:
| AttributeError: if the layer is connected to
| more than one incoming layers.
|
| trainable_variables
| Sequence of trainable variables owned by this module and its submodules.
|
| Note: this method uses reflection to find variables on the current instan
ce
| and submodules. For performance reasons you may wish to cache the result
| of calling this method if you don't expect the return value to change.
|
| Returns:
| A sequence of variables for the current module (sorted by attribute
| name) followed by variables from all submodules recursively (breadth
| first).
|
| updates
|
| variable_dtype
| Alias of `Layer.dtype`, the dtype of the weights.
|
| variables
| Returns the list of all layer variables/weights.
|
| Alias of `self.weights`.
|
| Note: This will not track the weights of nested `tf.Modules` that are
| not themselves Keras layers.
|
| Returns:
| A list of variables.
|
| -----
| Data descriptors inherited from keras.engine.base_layer.Layer:
|
| activity_regularizer
| Optional regularizer function for the output of this layer.
|
| stateful
|
| supports_masking
| Whether this layer supports computing a mask using `compute_mask`.
|
| trainable

```



-----  
 Class methods inherited from tensorflow.python.module.module.Module:

`with_name_scope(method)` from `builtins.type`

Decorator to automatically enter the module name scope.

```
>>> class MyModule(tf.Module):
... @tf.Module.with_name_scope
... def __call__(self, x):
... if not hasattr(self, 'w'):
... self.w = tf.Variable(tf.random.normal([x.shape[1], 3]))
... return tf.matmul(x, self.w)
```

Using the above module would produce ``tf.Variable`s` and ``tf.Tensor`s` whos

names included the module name:

```
>>> mod = MyModule()
>>> mod(tf.ones([1, 2]))
<tf.Tensor: shape=(1, 3), dtype=float32, numpy=..., dtype=float32>
>>> mod.w
<tf.Variable 'my_module/Variable:0' shape=(2, 3) dtype=float32,
numpy=..., dtype=float32>
```

Args:

method: The method to wrap.

Returns:

The original method wrapped such that it enters the module's name scop

-----  
 Readonly properties inherited from tensorflow.python.module.module.Module:

`name_scope`

Returns a ``tf.name_scope`` instance for this class.

`submodules`

Sequence of all sub-modules.

Submodules are modules which are properties of this module, or found as properties of modules which are properties of this module (and so on).

```
>>> a = tf.Module()
>>> b = tf.Module()
>>> c = tf.Module()
>>> a.b = b
>>> b.c = c
>>> list(a.submodules) == [b, c]
True
>>> list(b.submodules) == [c]
True
>>> list(c.submodules) == []
True
```

Returns:

A sequence of all submodules.

-----  
 Data descriptors inherited from tensorflow.python.trackable.base.Trackable:

```
|
| __dict__
| dictionary for instance variables (if defined)
|
| __weakref__
| list of weak references to the object (if defined)
```

```
In [44]: model = Sequential([Dense(4, activation='relu'), Dense(4, activation='relu'), De
```

```
In [45]: model.compile(loss='mse')
```

```
In [46]: model.fit(x=X_train, y=y_train, epochs=250)
```

```
Epoch 1/250
22/22 [=====] - 1s 3ms/step - loss: 256669.6875
Epoch 2/250
22/22 [=====] - 0s 4ms/step - loss: 256581.6406
Epoch 3/250
22/22 [=====] - 0s 4ms/step - loss: 256494.8594
Epoch 4/250
22/22 [=====] - 0s 4ms/step - loss: 256400.7031
Epoch 5/250
22/22 [=====] - 0s 4ms/step - loss: 256296.0938
Epoch 6/250
22/22 [=====] - 0s 4ms/step - loss: 256178.1719
Epoch 7/250
22/22 [=====] - 0s 4ms/step - loss: 256045.4844
Epoch 8/250
22/22 [=====] - 0s 4ms/step - loss: 255892.7344
Epoch 9/250
22/22 [=====] - 0s 4ms/step - loss: 255719.9375
Epoch 10/250
22/22 [=====] - 0s 4ms/step - loss: 255525.1719
Epoch 11/250
22/22 [=====] - 0s 4ms/step - loss: 255307.2031
Epoch 12/250
22/22 [=====] - 0s 4ms/step - loss: 255066.0781
Epoch 13/250
22/22 [=====] - 0s 4ms/step - loss: 254801.5469
Epoch 14/250
22/22 [=====] - 0s 4ms/step - loss: 254512.7031
Epoch 15/250
22/22 [=====] - 0s 4ms/step - loss: 254196.9844
Epoch 16/250
22/22 [=====] - 0s 5ms/step - loss: 253856.0938
Epoch 17/250
22/22 [=====] - 0s 4ms/step - loss: 253485.2344
Epoch 18/250
22/22 [=====] - 0s 4ms/step - loss: 253083.0156
Epoch 19/250
22/22 [=====] - 0s 4ms/step - loss: 252649.5156
Epoch 20/250
22/22 [=====] - 0s 4ms/step - loss: 252184.2969
Epoch 21/250
22/22 [=====] - 0s 3ms/step - loss: 251681.9844
Epoch 22/250
22/22 [=====] - 0s 4ms/step - loss: 251143.2031
Epoch 23/250
22/22 [=====] - 0s 3ms/step - loss: 250567.3438
Epoch 24/250
22/22 [=====] - 0s 4ms/step - loss: 249948.1094
Epoch 25/250
22/22 [=====] - 0s 4ms/step - loss: 249287.4219
Epoch 26/250
22/22 [=====] - 0s 4ms/step - loss: 248581.7188
Epoch 27/250
22/22 [=====] - 0s 4ms/step - loss: 247829.1719
Epoch 28/250
22/22 [=====] - 0s 3ms/step - loss: 247030.4531
Epoch 29/250
22/22 [=====] - 0s 6ms/step - loss: 246181.9844
Epoch 30/250
22/22 [=====] - 0s 4ms/step - loss: 245276.7344
```

```
Epoch 31/250
22/22 [=====] - 0s 3ms/step - loss: 244323.3125
Epoch 32/250
22/22 [=====] - 0s 4ms/step - loss: 243310.9688
Epoch 33/250
22/22 [=====] - 0s 4ms/step - loss: 242242.1719
Epoch 34/250
22/22 [=====] - 0s 3ms/step - loss: 241113.0000
Epoch 35/250
22/22 [=====] - 0s 3ms/step - loss: 239920.5469
Epoch 36/250
22/22 [=====] - 0s 4ms/step - loss: 238674.9531
Epoch 37/250
22/22 [=====] - 0s 4ms/step - loss: 237355.2031
Epoch 38/250
22/22 [=====] - 0s 3ms/step - loss: 235961.9219
Epoch 39/250
22/22 [=====] - 0s 3ms/step - loss: 234508.8906
Epoch 40/250
22/22 [=====] - 0s 3ms/step - loss: 232985.7344
Epoch 41/250
22/22 [=====] - 0s 3ms/step - loss: 231390.1094
Epoch 42/250
22/22 [=====] - 0s 4ms/step - loss: 229721.6406
Epoch 43/250
22/22 [=====] - 0s 4ms/step - loss: 227970.1094
Epoch 44/250
22/22 [=====] - 0s 4ms/step - loss: 226135.7656
Epoch 45/250
22/22 [=====] - 0s 4ms/step - loss: 224232.6094
Epoch 46/250
22/22 [=====] - 0s 5ms/step - loss: 222253.7188
Epoch 47/250
22/22 [=====] - 0s 5ms/step - loss: 220172.6094
Epoch 48/250
22/22 [=====] - 0s 5ms/step - loss: 218014.5781
Epoch 49/250
22/22 [=====] - 0s 5ms/step - loss: 215776.2031
Epoch 50/250
22/22 [=====] - 0s 4ms/step - loss: 213457.5312
Epoch 51/250
22/22 [=====] - 0s 5ms/step - loss: 211037.1406
Epoch 52/250
22/22 [=====] - 0s 4ms/step - loss: 208524.8906
Epoch 53/250
22/22 [=====] - 0s 4ms/step - loss: 205933.3906
Epoch 54/250
22/22 [=====] - 0s 4ms/step - loss: 203253.5938
Epoch 55/250
22/22 [=====] - 0s 4ms/step - loss: 200471.0000
Epoch 56/250
22/22 [=====] - 0s 3ms/step - loss: 197599.3594
Epoch 57/250
22/22 [=====] - 0s 5ms/step - loss: 194620.5938
Epoch 58/250
22/22 [=====] - 0s 3ms/step - loss: 191569.8906
Epoch 59/250
22/22 [=====] - 0s 3ms/step - loss: 188417.8281
Epoch 60/250
22/22 [=====] - 0s 3ms/step - loss: 185167.3750
```

```
Epoch 61/250
22/22 [=====] - 0s 3ms/step - loss: 181835.9062
Epoch 62/250
22/22 [=====] - 0s 3ms/step - loss: 178401.3750
Epoch 63/250
22/22 [=====] - 0s 4ms/step - loss: 174880.5625
Epoch 64/250
22/22 [=====] - 0s 6ms/step - loss: 171265.0625
Epoch 65/250
22/22 [=====] - 0s 4ms/step - loss: 167570.4062
Epoch 66/250
22/22 [=====] - 0s 4ms/step - loss: 163775.4375
Epoch 67/250
22/22 [=====] - 0s 3ms/step - loss: 159899.5938
Epoch 68/250
22/22 [=====] - 0s 3ms/step - loss: 155925.2656
Epoch 69/250
22/22 [=====] - 0s 3ms/step - loss: 151893.5781
Epoch 70/250
22/22 [=====] - 0s 4ms/step - loss: 147772.7031
Epoch 71/250
22/22 [=====] - 0s 3ms/step - loss: 143573.9531
Epoch 72/250
22/22 [=====] - 0s 3ms/step - loss: 139290.9219
Epoch 73/250
22/22 [=====] - 0s 4ms/step - loss: 134926.9062
Epoch 74/250
22/22 [=====] - 0s 3ms/step - loss: 130508.2969
Epoch 75/250
22/22 [=====] - 0s 4ms/step - loss: 126040.1484
Epoch 76/250
22/22 [=====] - 0s 12ms/step - loss: 121489.2891
Epoch 77/250
22/22 [=====] - 0s 7ms/step - loss: 116908.6406
Epoch 78/250
22/22 [=====] - 0s 4ms/step - loss: 112268.8359
Epoch 79/250
22/22 [=====] - 0s 4ms/step - loss: 107586.0078
Epoch 80/250
22/22 [=====] - 0s 3ms/step - loss: 102905.5859
Epoch 81/250
22/22 [=====] - 0s 3ms/step - loss: 98153.2812
Epoch 82/250
22/22 [=====] - 0s 5ms/step - loss: 93384.5703
Epoch 83/250
22/22 [=====] - 0s 4ms/step - loss: 88642.5469
Epoch 84/250
22/22 [=====] - 0s 4ms/step - loss: 83875.5469
Epoch 85/250
22/22 [=====] - 0s 3ms/step - loss: 79104.0781
Epoch 86/250
22/22 [=====] - 0s 3ms/step - loss: 74388.1406
Epoch 87/250
22/22 [=====] - 0s 4ms/step - loss: 69688.1641
Epoch 88/250
22/22 [=====] - 0s 4ms/step - loss: 64994.6562
Epoch 89/250
22/22 [=====] - 0s 3ms/step - loss: 60384.3477
Epoch 90/250
22/22 [=====] - 0s 4ms/step - loss: 55838.3672
```

```
Epoch 91/250
22/22 [=====] - 0s 6ms/step - loss: 51377.1953
Epoch 92/250
22/22 [=====] - 0s 4ms/step - loss: 46998.4648
Epoch 93/250
22/22 [=====] - 0s 3ms/step - loss: 42740.6445
Epoch 94/250
22/22 [=====] - 0s 4ms/step - loss: 38587.2734
Epoch 95/250
22/22 [=====] - 0s 3ms/step - loss: 34578.6484
Epoch 96/250
22/22 [=====] - 0s 3ms/step - loss: 30717.8086
Epoch 97/250
22/22 [=====] - 0s 3ms/step - loss: 27046.0137
Epoch 98/250
22/22 [=====] - 0s 3ms/step - loss: 23593.9805
Epoch 99/250
22/22 [=====] - 0s 3ms/step - loss: 20359.0918
Epoch 100/250
22/22 [=====] - 0s 4ms/step - loss: 17321.1543
Epoch 101/250
22/22 [=====] - 0s 4ms/step - loss: 14556.8359
Epoch 102/250
22/22 [=====] - 0s 3ms/step - loss: 12057.3545
Epoch 103/250
22/22 [=====] - 0s 3ms/step - loss: 9855.9365
Epoch 104/250
22/22 [=====] - 0s 3ms/step - loss: 7980.1455
Epoch 105/250
22/22 [=====] - 0s 4ms/step - loss: 6416.1924
Epoch 106/250
22/22 [=====] - 0s 3ms/step - loss: 5202.8813
Epoch 107/250
22/22 [=====] - 0s 3ms/step - loss: 4367.7812
Epoch 108/250
22/22 [=====] - 0s 3ms/step - loss: 3844.0901
Epoch 109/250
22/22 [=====] - 0s 3ms/step - loss: 3610.0786
Epoch 110/250
22/22 [=====] - 0s 3ms/step - loss: 3531.3564
Epoch 111/250
22/22 [=====] - 0s 3ms/step - loss: 3496.9631
Epoch 112/250
22/22 [=====] - 0s 3ms/step - loss: 3468.5608
Epoch 113/250
22/22 [=====] - 0s 3ms/step - loss: 3434.3894
Epoch 114/250
22/22 [=====] - 0s 3ms/step - loss: 3399.5039
Epoch 115/250
22/22 [=====] - 0s 3ms/step - loss: 3366.1819
Epoch 116/250
22/22 [=====] - 0s 3ms/step - loss: 3335.4500
Epoch 117/250
22/22 [=====] - 0s 3ms/step - loss: 3305.3606
Epoch 118/250
22/22 [=====] - 0s 3ms/step - loss: 3273.0967
Epoch 119/250
22/22 [=====] - 0s 3ms/step - loss: 3242.9424
Epoch 120/250
22/22 [=====] - 0s 3ms/step - loss: 3206.9714
```

```
Epoch 121/250
22/22 [=====] - 0s 3ms/step - loss: 3179.1804
Epoch 122/250
22/22 [=====] - 0s 3ms/step - loss: 3146.8735
Epoch 123/250
22/22 [=====] - 0s 3ms/step - loss: 3118.6758
Epoch 124/250
22/22 [=====] - 0s 3ms/step - loss: 3084.9307
Epoch 125/250
22/22 [=====] - 0s 3ms/step - loss: 3055.4043
Epoch 126/250
22/22 [=====] - 0s 3ms/step - loss: 3022.8354
Epoch 127/250
22/22 [=====] - 0s 3ms/step - loss: 2991.0105
Epoch 128/250
22/22 [=====] - 0s 3ms/step - loss: 2958.6587
Epoch 129/250
22/22 [=====] - 0s 3ms/step - loss: 2928.7686
Epoch 130/250
22/22 [=====] - 0s 3ms/step - loss: 2894.1025
Epoch 131/250
22/22 [=====] - 0s 6ms/step - loss: 2862.6543
Epoch 132/250
22/22 [=====] - 0s 3ms/step - loss: 2827.6001
Epoch 133/250
22/22 [=====] - 0s 3ms/step - loss: 2798.2551
Epoch 134/250
22/22 [=====] - 0s 4ms/step - loss: 2768.4961
Epoch 135/250
22/22 [=====] - 0s 5ms/step - loss: 2734.8162
Epoch 136/250
22/22 [=====] - 0s 3ms/step - loss: 2711.1099
Epoch 137/250
22/22 [=====] - 0s 3ms/step - loss: 2681.1011
Epoch 138/250
22/22 [=====] - 0s 5ms/step - loss: 2644.6653
Epoch 139/250
22/22 [=====] - 0s 4ms/step - loss: 2617.4651
Epoch 140/250
22/22 [=====] - 0s 4ms/step - loss: 2589.4026
Epoch 141/250
22/22 [=====] - 0s 4ms/step - loss: 2558.7117
Epoch 142/250
22/22 [=====] - 0s 4ms/step - loss: 2531.3167
Epoch 143/250
22/22 [=====] - 0s 4ms/step - loss: 2504.3027
Epoch 144/250
22/22 [=====] - 0s 4ms/step - loss: 2472.5715
Epoch 145/250
22/22 [=====] - 0s 3ms/step - loss: 2447.0303
Epoch 146/250
22/22 [=====] - 0s 5ms/step - loss: 2417.0134
Epoch 147/250
22/22 [=====] - 0s 5ms/step - loss: 2388.3555
Epoch 148/250
22/22 [=====] - 0s 8ms/step - loss: 2357.0532
Epoch 149/250
22/22 [=====] - 0s 4ms/step - loss: 2326.2117
Epoch 150/250
22/22 [=====] - 0s 5ms/step - loss: 2294.7671
```

```
Epoch 151/250
22/22 [=====] - 0s 8ms/step - loss: 2262.1885
Epoch 152/250
22/22 [=====] - 0s 5ms/step - loss: 2238.1206
Epoch 153/250
22/22 [=====] - 0s 6ms/step - loss: 2208.6887
Epoch 154/250
22/22 [=====] - 0s 4ms/step - loss: 2179.6489
Epoch 155/250
22/22 [=====] - 0s 4ms/step - loss: 2151.7029
Epoch 156/250
22/22 [=====] - 0s 3ms/step - loss: 2128.5547
Epoch 157/250
22/22 [=====] - 0s 3ms/step - loss: 2102.6523
Epoch 158/250
22/22 [=====] - 0s 4ms/step - loss: 2073.4790
Epoch 159/250
22/22 [=====] - 0s 2ms/step - loss: 2047.2493
Epoch 160/250
22/22 [=====] - 0s 2ms/step - loss: 2021.5590
Epoch 161/250
22/22 [=====] - 0s 2ms/step - loss: 1997.5389
Epoch 162/250
22/22 [=====] - 0s 2ms/step - loss: 1966.7186
Epoch 163/250
22/22 [=====] - 0s 3ms/step - loss: 1938.3456
Epoch 164/250
22/22 [=====] - 0s 2ms/step - loss: 1911.6132
Epoch 165/250
22/22 [=====] - 0s 2ms/step - loss: 1883.1777
Epoch 166/250
22/22 [=====] - 0s 2ms/step - loss: 1855.9683
Epoch 167/250
22/22 [=====] - 0s 2ms/step - loss: 1830.6466
Epoch 168/250
22/22 [=====] - 0s 2ms/step - loss: 1803.8029
Epoch 169/250
22/22 [=====] - 0s 2ms/step - loss: 1780.0208
Epoch 170/250
22/22 [=====] - 0s 2ms/step - loss: 1753.9244
Epoch 171/250
22/22 [=====] - 0s 2ms/step - loss: 1728.4247
Epoch 172/250
22/22 [=====] - 0s 3ms/step - loss: 1701.5580
Epoch 173/250
22/22 [=====] - 0s 2ms/step - loss: 1674.8873
Epoch 174/250
22/22 [=====] - 0s 2ms/step - loss: 1647.7418
Epoch 175/250
22/22 [=====] - 0s 2ms/step - loss: 1621.2034
Epoch 176/250
22/22 [=====] - 0s 2ms/step - loss: 1595.6647
Epoch 177/250
22/22 [=====] - 0s 2ms/step - loss: 1566.8955
Epoch 178/250
22/22 [=====] - 0s 2ms/step - loss: 1542.1577
Epoch 179/250
22/22 [=====] - 0s 2ms/step - loss: 1510.2678
Epoch 180/250
22/22 [=====] - 0s 2ms/step - loss: 1487.2109
```



```
Epoch 181/250
22/22 [=====] - 0s 2ms/step - loss: 1459.6084
Epoch 182/250
22/22 [=====] - 0s 1ms/step - loss: 1432.4683
Epoch 183/250
22/22 [=====] - 0s 2ms/step - loss: 1410.1958
Epoch 184/250
22/22 [=====] - 0s 2ms/step - loss: 1382.9072
Epoch 185/250
22/22 [=====] - 0s 2ms/step - loss: 1360.2345
Epoch 186/250
22/22 [=====] - 0s 2ms/step - loss: 1335.9716
Epoch 187/250
22/22 [=====] - 0s 2ms/step - loss: 1311.9059
Epoch 188/250
22/22 [=====] - 0s 2ms/step - loss: 1285.9893
Epoch 189/250
22/22 [=====] - 0s 2ms/step - loss: 1261.5288
Epoch 190/250
22/22 [=====] - 0s 2ms/step - loss: 1236.7725
Epoch 191/250
22/22 [=====] - 0s 2ms/step - loss: 1212.1093
Epoch 192/250
22/22 [=====] - 0s 2ms/step - loss: 1186.6194
Epoch 193/250
22/22 [=====] - 0s 2ms/step - loss: 1164.6515
Epoch 194/250
22/22 [=====] - 0s 4ms/step - loss: 1138.6422
Epoch 195/250
22/22 [=====] - 0s 2ms/step - loss: 1117.6681
Epoch 196/250
22/22 [=====] - 0s 2ms/step - loss: 1095.6549
Epoch 197/250
22/22 [=====] - 0s 2ms/step - loss: 1072.4618
Epoch 198/250
22/22 [=====] - 0s 2ms/step - loss: 1051.4398
Epoch 199/250
22/22 [=====] - 0s 2ms/step - loss: 1030.2701
Epoch 200/250
22/22 [=====] - 0s 2ms/step - loss: 1009.5002
Epoch 201/250
22/22 [=====] - 0s 2ms/step - loss: 988.6779
Epoch 202/250
22/22 [=====] - 0s 2ms/step - loss: 965.1537
Epoch 203/250
22/22 [=====] - 0s 2ms/step - loss: 942.3541
Epoch 204/250
22/22 [=====] - 0s 2ms/step - loss: 918.0474
Epoch 205/250
22/22 [=====] - 0s 2ms/step - loss: 894.1251
Epoch 206/250
22/22 [=====] - 0s 2ms/step - loss: 871.5606
Epoch 207/250
22/22 [=====] - 0s 2ms/step - loss: 846.8628
Epoch 208/250
22/22 [=====] - 0s 2ms/step - loss: 827.7665
Epoch 209/250
22/22 [=====] - 0s 2ms/step - loss: 805.4969
Epoch 210/250
22/22 [=====] - 0s 2ms/step - loss: 784.8796
```

```
Epoch 211/250
22/22 [=====] - 0s 2ms/step - loss: 765.1651
Epoch 212/250
22/22 [=====] - 0s 2ms/step - loss: 744.1163
Epoch 213/250
22/22 [=====] - 0s 2ms/step - loss: 724.2258
Epoch 214/250
22/22 [=====] - 0s 4ms/step - loss: 703.8245
Epoch 215/250
22/22 [=====] - 0s 2ms/step - loss: 682.2206
Epoch 216/250
22/22 [=====] - 0s 2ms/step - loss: 663.8043
Epoch 217/250
22/22 [=====] - 0s 2ms/step - loss: 644.1346
Epoch 218/250
22/22 [=====] - 0s 3ms/step - loss: 624.1998
Epoch 219/250
22/22 [=====] - 0s 2ms/step - loss: 602.5823
Epoch 220/250
22/22 [=====] - 0s 2ms/step - loss: 582.3586
Epoch 221/250
22/22 [=====] - 0s 2ms/step - loss: 563.2145
Epoch 222/250
22/22 [=====] - 0s 2ms/step - loss: 544.9669
Epoch 223/250
22/22 [=====] - 0s 2ms/step - loss: 527.1445
Epoch 224/250
22/22 [=====] - 0s 2ms/step - loss: 508.8384
Epoch 225/250
22/22 [=====] - 0s 2ms/step - loss: 489.9565
Epoch 226/250
22/22 [=====] - 0s 2ms/step - loss: 469.9984
Epoch 227/250
22/22 [=====] - 0s 2ms/step - loss: 457.8904
Epoch 228/250
22/22 [=====] - 0s 2ms/step - loss: 439.4278
Epoch 229/250
22/22 [=====] - 0s 2ms/step - loss: 423.0347
Epoch 230/250
22/22 [=====] - 0s 2ms/step - loss: 404.9425
Epoch 231/250
22/22 [=====] - 0s 2ms/step - loss: 388.3860
Epoch 232/250
22/22 [=====] - 0s 4ms/step - loss: 371.7105
Epoch 233/250
22/22 [=====] - 0s 2ms/step - loss: 356.8831
Epoch 234/250
22/22 [=====] - 0s 2ms/step - loss: 341.6101
Epoch 235/250
22/22 [=====] - 0s 2ms/step - loss: 327.6884
Epoch 236/250
22/22 [=====] - 0s 2ms/step - loss: 312.1833
Epoch 237/250
22/22 [=====] - 0s 2ms/step - loss: 298.5418
Epoch 238/250
22/22 [=====] - 0s 2ms/step - loss: 284.2587
Epoch 239/250
22/22 [=====] - 0s 2ms/step - loss: 270.4089
Epoch 240/250
22/22 [=====] - 0s 2ms/step - loss: 258.5845
```

```

Epoch 241/250
22/22 [=====] - 0s 2ms/step - loss: 245.6975
Epoch 242/250
22/22 [=====] - 0s 2ms/step - loss: 233.6224
Epoch 243/250
22/22 [=====] - 0s 2ms/step - loss: 221.3953
Epoch 244/250
22/22 [=====] - 0s 2ms/step - loss: 210.0626
Epoch 245/250
22/22 [=====] - 0s 2ms/step - loss: 197.9543
Epoch 246/250
22/22 [=====] - 0s 2ms/step - loss: 185.6946
Epoch 247/250
22/22 [=====] - 0s 2ms/step - loss: 174.8311
Epoch 248/250
22/22 [=====] - 0s 2ms/step - loss: 164.2820
Epoch 249/250
22/22 [=====] - 0s 4ms/step - loss: 154.3989
Epoch 250/250
22/22 [=====] - 0s 2ms/step - loss: 144.5810

```

Out[46]: <keras.callbacks.History at 0x27b3191a670>

```
In [50]: loss_df = pd.DataFrame(model.history.history)
```

```
In [54]: loss_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 0 entries
Empty DataFrame

```

```
In [41]: #Evaluation of model; how well model is perform on the data never seen before
```

```
In [55]: model.evaluate(X_test, y_test, verbose=0)
```

Out[55]: 135.16517639160156

```
In [56]: model.evaluate(X_train, y_train, verbose=0)
```

Out[56]: 139.91188049316406

```
In [57]: test_predictions = model.predict(X_test)
```

```
10/10 [=====] - 0s 2ms/step
```

```
In [58]: test_predictions
```

```
Out[58]: array([[418.39847],
 [612.6908],
 [585.17664],
 [565.9918],
 [382.02237],
 [572.9095],
 [514.2187],
 [467.0733],
 [545.5685],
 [456.91183],
 [602.2751],
 [548.21515],
 [430.545],
 [421.13052],
 [639.3819],
 [447.4427],
 [511.84842],
 [643.42316],
 [646.7347],
 [561.8312],
 [354.13635],
 [453.4814],
 [397.39618],
 [394.8635],
 [561.8605],
 [600.14435],
 [531.2583],
 [438.675],
 [639.9403],
 [426.9883],
 [452.39566],
 [489.4803],
 [448.03287],
 [663.7151],
 [437.00095],
 [428.6107],
 [506.50925],
 [547.8075],
 [509.8763],
 [407.70612],
 [608.71405],
 [429.06094],
 [596.0527],
 [455.98355],
 [504.44662],
 [577.5778],
 [654.72705],
 [494.92734],
 [342.08957],
 [489.55356],
 [518.3941],
 [398.39413],
 [541.4695],
 [421.40778],
 [629.20514],
 [494.41794],
 [615.09357],
 [617.41705],
 [455.7144],
 [488.155]])
```

[494.10474],  
[480.56625],  
[665.77203],  
[415.96927],  
[681.8936 ],  
[579.096 ],  
[576.91376],  
[537.2094 ],  
[488.88852],  
[516.5376 ],  
[379.39578],  
[538.98236],  
[565.1118 ],  
[529.9635 ],  
[461.85483],  
[530.96765],  
[510.27188],  
[451.22086],  
[540.95636],  
[629.2926 ],  
[470.64835],  
[563.05914],  
[673.83984],  
[465.36325],  
[687.60034],  
[478.23367],  
[418.00098],  
[578.314 ],  
[447.9073 ],  
[493.599 ],  
[607.2725 ],  
[449.74042],  
[461.9061 ],  
[444.73264],  
[507.71973],  
[598.77264],  
[343.78043],  
[446.44943],  
[535.524 ],  
[518.44086],  
[595.007 ],  
[525.00946],  
[356.10358],  
[571.86847],  
[439.96378],  
[560.39636],  
[514.8572 ],  
[405.1473 ],  
[559.8668 ],  
[464.2247 ],  
[456.2629 ],  
[628.3588 ],  
[525.43787],  
[548.75323],  
[429.35208],  
[483.43463],  
[579.7542 ],  
[650.8828 ],  
[680.39874],  
[642.7178 ],

[556.6105 ],  
[506.4194 ],  
[403.84946],  
[307.50143],  
[485.01193],  
[607.5755 ],  
[391.24716],  
[513.7102 ],  
[511.06427],  
[496.17917],  
[485.40298],  
[433.56375],  
[496.48868],  
[477.48938],  
[593.5266 ],  
[570.07336],  
[425.3471 ],  
[619.324 ],  
[473.44077],  
[560.78046],  
[418.6245 ],  
[531.1808 ],  
[567.2133 ],  
[374.80664],  
[546.5545 ],  
[594.6292 ],  
[399.2957 ],  
[539.5378 ],  
[558.53015],  
[461.73172],  
[620.5473 ],  
[387.56235],  
[478.8971 ],  
[527.1921 ],  
[389.6014 ],  
[467.23172],  
[445.03046],  
[501.1993 ],  
[366.8724 ],  
[410.26016],  
[596.3274 ],  
[509.0742 ],  
[474.50082],  
[494.9113 ],  
[535.5377 ],  
[363.9023 ],  
[516.0079 ],  
[281.4089 ],  
[506.40277],  
[537.6828 ],  
[492.96793],  
[478.91653],  
[407.23288],  
[429.84964],  
[546.37256],  
[480.0657 ],  
[572.9029 ],  
[492.37854],  
[590.9702 ],  
[545.38715],

[540.6818 ],  
[503.37418],  
[632.12555],  
[556.1739 ],  
[573.93945],  
[454.4816 ],  
[428.63748],  
[430.34937],  
[565.1976 ],  
[599.1899 ],  
[448.15427],  
[492.45468],  
[579.6175 ],  
[524.42303],  
[375.34586],  
[631.5297 ],  
[526.87463],  
[357.8071 ],  
[496.06485],  
[423.91006],  
[597.92523],  
[367.28464],  
[523.80304],  
[418.46045],  
[285.68088],  
[521.1608 ],  
[362.98608],  
[378.54532],  
[572.0422 ],  
[429.75912],  
[547.4466 ],  
[520.89264],  
[509.64145],  
[347.1237 ],  
[415.82785],  
[592.862 ],  
[606.8047 ],  
[592.5543 ],  
[563.13446],  
[478.00607],  
[467.5994 ],  
[509.9711 ],  
[455.4181 ],  
[514.17737],  
[503.1556 ],  
[414.43945],  
[596.5465 ],  
[287.09836],  
[617.40094],  
[584.25006],  
[349.77612],  
[483.4514 ],  
[587.9591 ],  
[394.7569 ],  
[467.08432],  
[348.02017],  
[520.4397 ],  
[423.72888],  
[553.3029 ],  
[629.54816],

```
[535.23456],
[506.22397],
[623.9798],
[516.91815],
[530.92285],
[518.5015],
[465.1708],
[508.99954],
[468.5614],
[586.03955],
[471.9312],
[436.9844],
[538.67346],
[496.90576],
[660.0891],
[389.309],
[548.74335],
[572.2406],
[442.9645],
[540.42957],
[580.65283],
[574.46295],
[699.1895],
[441.23053],
[413.08182],
[337.7723],
[456.0394],
[404.7378],
[540.3757],
[522.61505],
[561.8919],
[456.5273],
[533.8483],
[397.80563],
[504.33624],
[623.902],
[501.2042],
[562.8967],
[475.38162],
[302.90704],
[517.3773],
[610.40784],
[369.5964],
[457.173],
[501.33182],
[540.87195],
[603.0571],
[401.57938],
[456.0519],
[487.51392],
[589.7326],
[500.72278],
[343.8138],
[551.70764],
[453.5452],
[527.69965],
[513.18054],
[600.54816],
[429.69885],
[422.62534]], dtype=float32)
```



```
In [59]: test_predictions = pd.Series(test_predictions.reshape(300,))
```

```
In [60]: test_predictions
```

```
Out[60]: 0 418.398468
 1 612.690796
 2 585.176636
 3 565.991821
 4 382.022369
 ...
 295 527.699646
 296 513.180542
 297 600.548157
 298 429.698853
 299 422.625336
 Length: 300, dtype: float32
```

```
In [61]: pred_df = pd.DataFrame(y_test, columns=['Test True Y'])
```

```
In [62]: pred_df = pd.concat([pred_df, test_predictions], axis=1)
```

```
In [63]: pred_df.columns = ['Test True Y', 'Model Predictions']
```

```
In [64]: pred_df
```


```
Out[64]:
```

	Test True Y	Model Predictions
0	402.296319	418.398468
1	624.156198	612.690796
2	582.455066	585.176636
3	578.588606	565.991821
4	371.224104	382.022369
...	...	...
295	525.704657	527.699646
296	502.909473	513.180542
297	612.727910	600.548157
298	417.569725	429.698853
299	410.538250	422.625336

300 rows × 2 columns

```
In [65]: sns.scatterplot(x='Test True Y', y='Model Predictions', data=pred_df)
```

```
Out[65]: <Axes: xlabel='Test True Y', ylabel='Model Predictions'>
```



In [66]: *#The above chart represents the model is working very well*

In [67]: `from sklearn.metrics import mean_absolute_error, mean_squared_error`

In [68]: `mean_absolute_error(pred_df['Test True Y'], pred_df['Model Predictions'])`

Out[68]: 9.319370338603512

In [69]: `df.describe()`

Out[69]:

	price	feature1	feature2
<b>count</b>	1000.000000	1000.000000	1000.000000
<b>mean</b>	498.673029	1000.014171	999.979847
<b>std</b>	93.785431	0.974018	0.948330
<b>min</b>	223.346793	997.058347	996.995651
<b>25%</b>	433.025732	999.332068	999.316106
<b>50%</b>	502.382117	1000.009915	1000.002243
<b>75%</b>	564.921588	1000.637580	1000.645380
<b>max</b>	774.407854	1003.207934	1002.666308

In [72]: *#Root mean Square error*  
`mean_squared_error(pred_df['Test True Y'], pred_df['Model Predictions'])*0.5`

Out[72]: 11.626055334189301

```
In [73]: new_gem = [[998, 1000]]
```

```
In [75]: new_gem = scaler.transform(new_gem)
```

```
In [76]: model.predict(new_gem)
```

```
1/1 [=====] - 0s 33ms/step
```

```
Out[76]: array([[429.40448]], dtype=float32)
```

```
In [77]: from tensorflow.keras.models import load_model
```

```
In [78]: model.save('my_gem_model.h5')
```

```
In [79]: later_model = load_model('my_gem_model.h5')
```

```
In [80]: later_model.predict(new_gem)
```

```
1/1 [=====] - 0s 198ms/step
```

```
Out[80]: array([[429.40448]], dtype=float32)
```

```
In []:
```